

CAFFE TUTORIAL



Maximally accurate

Maximally specific

espresso

2.23192

coffee

2.19914

beverage

1.93214

liquid

1.89367

fluid

1.85519

Brewing Deep Networks With Caffe

ROHIT GIRDHAR

Many slides from Xinlei Chen (16-824 tutorial), Caffe CVPR'15 tutorial


Overview

- Motivation and comparisons
- Training/Finetuning a simple model
- Deep dive into Caffe!

! this->tutorial


- What is Deep Learning?
- Why Deep Learning?
 - The Unreasonable Effectiveness of Deep Features
- History of Deep Learning.

CNNs 1989



LeNet: a layered model composed of convolution and subsampling operations followed by a holistic representation and ultimately a classifier for handwritten digits. [LeNet]

CNNs 2012



AlexNet: a layered model composed of convolution, subsampling, and further operations followed by a holistic representation and all-in-all a landmark classifier on ILSVRC12. [AlexNet]
+ data, + gpu, + non-saturating nonlinearity, + regularization

Other Frameworks

- Torch7
 - NYU
 - scientific computing framework in Lua
 - supported by Facebook
- TensorFlow
 - Google
 - Good for deploying
- Theano/Pylearn2
 - U. Montreal
 - scientific computing framework in Python
 - symbolic computation and automatic differentiation
- Cuda-Convnet2
 - Alex Krizhevsky
 - Very fast on state-of-the-art GPUs with Multi-GPU parallelism
 - C++ / CUDA library
- MatConvNet
 - Oxford U.
 - Deep Learning in MATLAB
- CXXNet
- Marvin



Framework Comparison

- More alike than different
 - All express deep models
 - All are open-source (contributions differ)
 - Most include scripting for hacking and prototyping
- No strict winners – experiment and choose the framework that best fits your work

Torch vs Caffe vs TensorFlow?

- Torch has more functionality built-in (more variety of layers etc.) and is in general more flexible
- However, more flexibility => writing more code! If you have a million images and want to train a mostly standard architecture, go with caffe!
- TensorFlow is best at deployment! Even works on mobile devices.

What is Caffe?

Open framework, models, and worked examples for deep learning

- 600+ citations, 150+ contributors, 7,000+ stars, 4,700+ forks, >1 pull request / day average
- focus has been vision, but branching out: sequences, reinforcement learning, speech + text

The screenshot shows the GitHub repository page for BVLC / caffe. At the top, the repository name is displayed with icons for Watch (1,202), Star (8,320), and Fork (4,728). Below this is a navigation bar with links for Code, Issues (333), Pull requests (188), Wiki, Pulse, and Graphs. The repository description reads: "Caffe: a fast open framework for deep learning. <http://caffe.berkeleyvision.org/>". A progress bar shows 3,523 commits, 6 branches, 11 releases, and 172 contributors. At the bottom, there are buttons for "New pull request", "New file", "Find file", "HTTPS" (with a dropdown), a link to "https://github.com/BVLC/c", and a "Download ZIP" button. A recent activity bar at the very bottom shows a merge pull request #3624 by jeffdonahue from dmikolaev/bvlc-print-gpu-names, with the latest commit b590f1d made 2 days ago.

BVLC / caffe

Watch 1,202 Star 8,320 Fork 4,728

Code Issues 333 Pull requests 188 Wiki Pulse Graphs

Caffe: a fast open framework for deep learning. <http://caffe.berkeleyvision.org/>

3,523 commits 6 branches 11 releases 172 contributors

Branch: master New pull request

New file Find file HTTPS <https://github.com/BVLC/c> Download ZIP

jeffdonahue Merge pull request #3624 from dmikolaev/bvlc-print-gpu-names Latest commit b590f1d 2 days ago

So what is Caffe?

- Pure C++ / CUDA architecture for deep learning
 - command line, Python, MATLAB interfaces
- Fast, well-tested code
- Tools, reference models, demos, and recipes
- Seamless switch between CPU and GPU
 - `Caffe::set_mode(Caffe::GPU);`



Prototype



Training



Deployment

All with essentially the same code!

Brewing by the Numbers...

- Speed with Krizhevsky's 2012 model:
 - 2 ms / image on K40 GPU
 - **<1 ms** inference with Caffe + cuDNN v2 on Titan X
 - **72 million** images / day with batched IO
 - 8-core CPU: ~20 ms/image
- **9k** lines of C++ code (20k with tests)
- <https://github.com/soumith/convnet-benchmarks>: A pretty reliable benchmark

Why Caffe? In one sip...

Expression: models + optimizations are plaintext schemas, not code.

Speed: for state-of-the-art models and massive data.

Modularity: to extend to new tasks and settings.

Openness: common code and reference models for reproducibility.

Community: joint discussion and development through BSD-2 licensing.

Caffe Tutorial

<http://caffe.berkeleyvision.org/tutorial/>


- **Nets, Layers, and Blobs**: the anatomy of a Caffe model.
- **Forward / Backward**: the essential computations of layered compositional models.
- **Loss**: the task to be learned is defined by the loss.
- **Solver**: the solver coordinates model optimization.
- **Layer Catalogue**: the layer is the fundamental unit of modeling and computation – Caffe's catalogue includes layers for state-of-the-art models.
- **Interfaces**: command line, Python, and MATLAB Caffe.
- **Data**: how to caffeinate data for model input.

For a closer look at a few details:


- **Caffeinated Convolution**: how Caffe computes convolutions.

Reference Models

AlexNet: ImageNet Classification



R-CNN: Regions with CNN features



Caffe offers the

- model definitions
- optimization settings
- pre-trained weights

so you can start right away.

The BVLC models are licensed for unrestricted use.

Open Model Collection


The Caffe [Model Zoo](#)

- open collection of deep models to share innovation
 - VGG ILSVRC14 + Devil models **in the zoo**
 - Network-in-Network / CCCP model **in the zoo**
 - MIT Places scene recognition model **in the zoo**
- help disseminate and reproduce research
- bundled tools for loading and publishing models

Share Your Models! with your citation + license of course


Architectures

DAGs
multi-input
multi-task



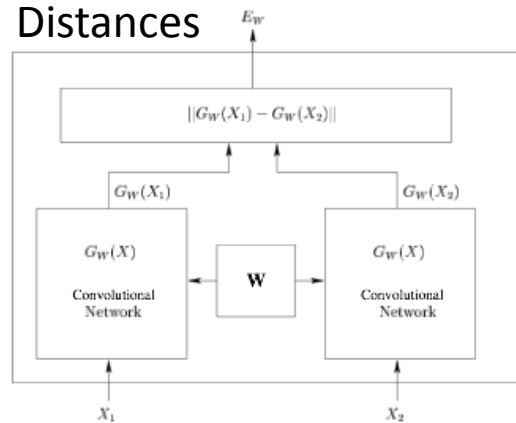
[Karpathy14]

Weight Sharing
Recurrent (RNNs)
Sequences



[Sutskever13]

Siamese Nets
Distances



[Chopra05]

Define your own model from our catalogue
of layers types and start learning.

Installation Hints

- We have already compiled the latest version of caffe (as on 5 Feb'16) on LateDays!
- However, you might want to customize and compile your own caffe (esp. if you want to create new layers)

Installation

- <http://caffe.berkeleyvision.org/installation.html>
- **CUDA, OPENCV**
- **BLAS** (Basic Linear Algebra Subprograms): operations like matrix multiplication, matrix addition, both implementation for CPU(cBLAS) and GPU(cuBLAS). provided by MKL(INTEL), ATLAS, openBLAS, etc.
- **Boost**: a c++ library. > Use some of its math functions and shared_pointer.
- **glog,gflags** provide logging & command line utilities. > Essential for debugging.
- **leveldb, lmdb**: database io for your program. > Need to know this for preparing your own data.
- **protobuf**: an efficient and flexible way to define data structure. > Need to know this for defining new layers.

TRAINING AND FINE-TUNING

Training: Step 1

Create a lenet_train.prototxt

Data

```
layer {
  name: "data"
  type: "Data"
  top: "data"
  top: "label"
  transform_param {
    scale: 0.00392156862745
  }
  data_param {
    source: "examples/mnist/mnist_train_lmdb"
    batch_size: 64
    backend: LMDB
  }
}
```

Layers

```
layer {
  name: "conv1"
  type: "Convolution"
  bottom: "data"
  top: "conv1"
  convolution_param {
    num_output: 20
    kernel_size: 5
    weight_filler {
      type: "xavier"
    }
  }
}
```

Loss

```
layer {
  name: "loss"
  type: "SoftmaxWithLoss"
  bottom: "ip2"
  bottom: "label"
  top: "loss"
}
```

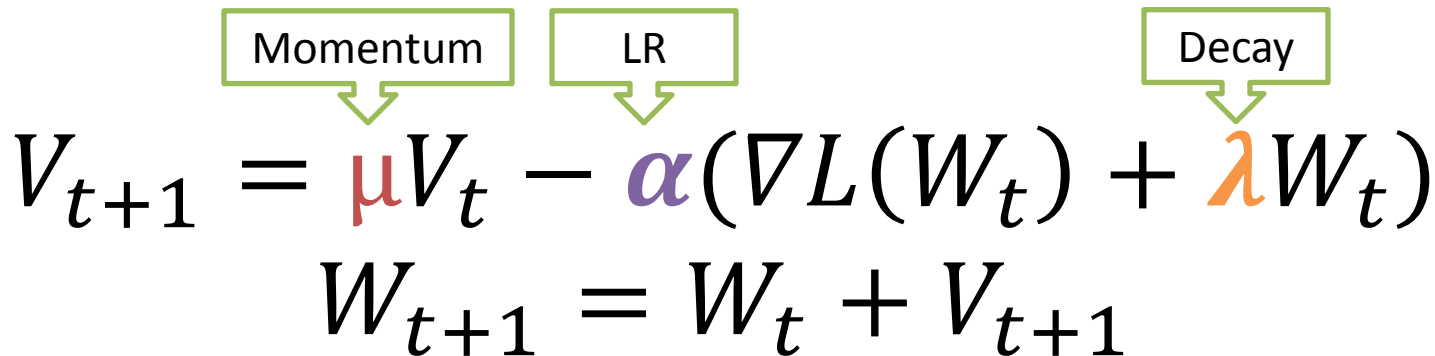
Training: Step 2

Create a `lenet_solver.prototxt`

```
train_net: "lenet_train.prototxt"  
base_lr: 0.01  
momentum: 0.9  
weight_decay: 0.0005  
max_iter: 10000  
snapshot_prefix: "lenet_snapshot"  
# ... and some other options ...
```

Training: Step 2

Some details on SGD parameters



The diagram illustrates the SGD update equations with three annotations in green boxes: 'Momentum' pointing to the red μ in the first equation, 'LR' pointing to the purple α in the first equation, and 'Decay' pointing to the orange λ in the first equation. The equations are:

$$V_{t+1} = \mu V_t - \alpha (\nabla L(W_t) + \lambda W_t)$$
$$W_{t+1} = W_t + V_{t+1}$$

Training: Step 3

Train!

```
$ build/tools/caffe train \  
  -solver lenet_solver.prototxt \  
  -gpu 0
```

Fine-tuning

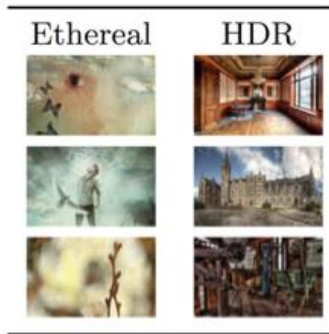
Transferring learned weights to kick-start models

- Take a pre-trained model and fine-tune to new tasks [DeCAF] [Zeiler-Fergus] [OverFeat]



dog bird invertebrate vehicle good, covering commodity building

Your Task



**Style
Recognition**



**Dogs vs.
Cats**
top 10 in
10 minutes

© kaggle.com

From ImageNet to Style

- Simply change a few lines in the layer definition

```
layers {  
  name: "data"  
  type: DATA  
  data_param {  
    source: "ilsvrc12_train_leveldb"  source: "style_leveldb"  
    mean_file: "../..data/ilsvrc12"  
    ...  
  }  
  ...  
}  
...  
layers {  
  name: "fc8"  name: "fc8-style" new name = new params  
  type: INNER_PRODUCT  
  blobs_lr: 1  
  blobs_lr: 2  
  weight_decay: 1  
  weight_decay: 0  
  inner_product_param {  
    num_output: 1000  num_output: 20  
    ...  
  }  
  ...  
}
```

Input: A different source

Last Layer: A different classifier

From ImageNet to Style

```
$ caffe train -solver models/finetune_flickr_style/solver.prototxt \  
             -gpu 0 \  
             -weights bvlc_reference_caffenet.caffemodel
```

Under the hood (loosely speaking):

```
net = new Caffe::Net(  
    "style_solver.prototxt");  
net.CopyTrainedNetFrom(  
    pretrained_model);  
solver.Solve(net);
```

Vintage HDR Melancholy Minimal




When to Fine-tune?

A good first step!

- More robust optimization – good initialization helps
- Needs less data
- Faster learning

State-of-the-art results in

- recognition
- detection
- segmentation



Fine-tuning Tricks

Learn the last layer first

- Caffe layers have local learning rates: `blobs_lr`
- Freeze all but the last layer for fast optimization and avoiding early divergence.
- Stop if good enough, or keep fine-tuning


Reduce the learning rate

- Drop the solver learning rate by 10x, 100x
- Preserve the initialization from pre-training and avoid thrashing

DEEPER INTO CAFFE

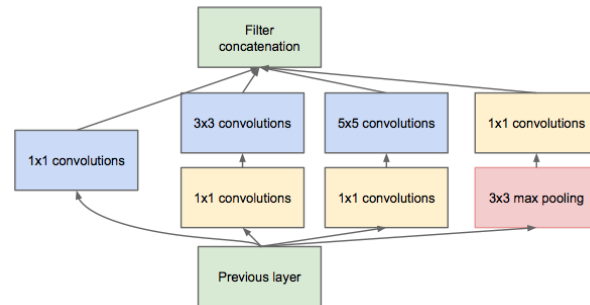

DAG

Many current deep models have linear structure

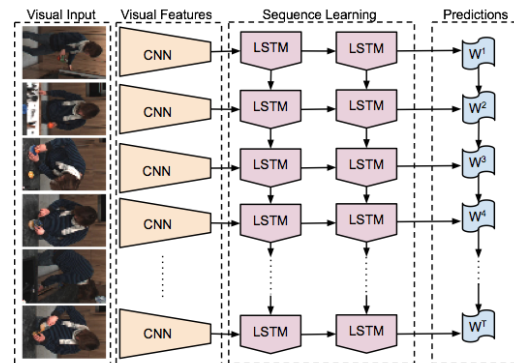


but Caffe nets can have any directed acyclic graph (DAG) structure.

Define bottoms and tops
and Caffe will connect the net.



GoogLeNet Inception Module




LRCN joint vision-sequence model

Net

- A network is a set of layers connected as a DAG:

```
name: "dummy-net"  
layers { name: "data" ...}  
layers { name: "conv" ...}  
layers { name: "pool" ...}  
... more layers ...  
layers { name: "loss" ...}
```

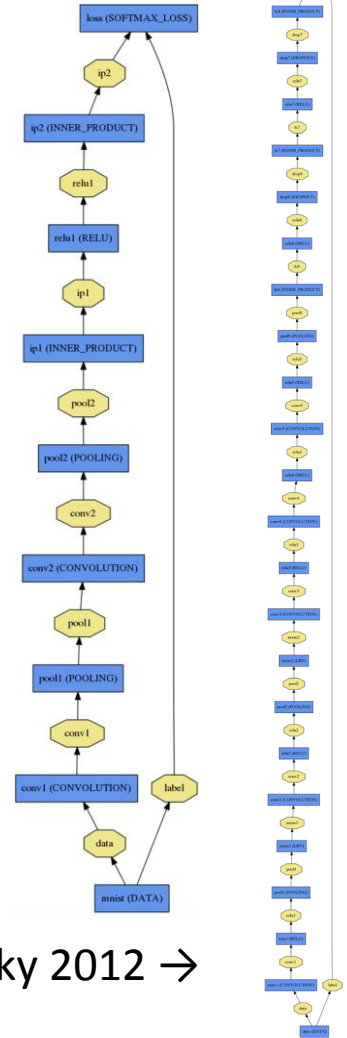
- Caffe creates and checks the net from the definition.
- Data and derivatives flow through the net as *blobs* – a an array interface



LogReg
↑

LeNet →

ImageNet, Krizhevsky 2012 →




Layer

```
name: "conv1"
type: "Convolution"
bottom: "data"
top: "conv1"
convolution_param {
  num_output: 20
  kernel_size: 5
  stride: 1
  weight_filler {
    type: "xavier"
  }
}
```

name, type, and the
connection
structure
(input blobs and
output blobs)

layer-specific
parameters



- Every layer type defines

- **Setup**
- **Forward**
- **Backward**

* Nets + Layers are defined by [protobuf](#) schema

Layer Protocol

Setup: run once for initialization.


Forward: make output given input.

Backward: make gradient of output

- w.r.t. bottom
- w.r.t. parameters (if needed)

Model Composition

The Net forward and backward passes are the composition the layers'.




[Layer Development Checklist](#)

Blob

N

Blobs are ~~4~~-D arrays for storing and communicating information.


- hold data, derivatives, and parameters
- lazily allocate memory
- shuttle between CPU and GPU



Data

Number x K Channel x Height x Width

256 x 3 x 227 x 227 for ImageNet train input



Parameter: Convolution Weight

N Output x K Input x Height x Width


96 x 3 x 11 x 11 for CaffeNet conv1



Parameter: Convolution Bias


96 x 1 x 1 x 1 for CaffeNet conv1

```
name: "conv1"  
type: "Convolution"  
bottom: "data"  
top: "conv1"  
... definition ...
```



Blob

Blobs provide a unified memory interface.



Reshape(num, channel, height, width)

- declare dimensions
- make *SyncedMem* -- but only lazily allocate

cpu_data(), mutable_cpu_data()

- host memory for CPU mode

gpu_data(), mutable_gpu_data()

- device memory for GPU mode

{cpu,gpu}_diff(), mutable_{cpu,gpu}_diff()

- derivative counterparts to data methods
- easy access to data + diff in forward / backward

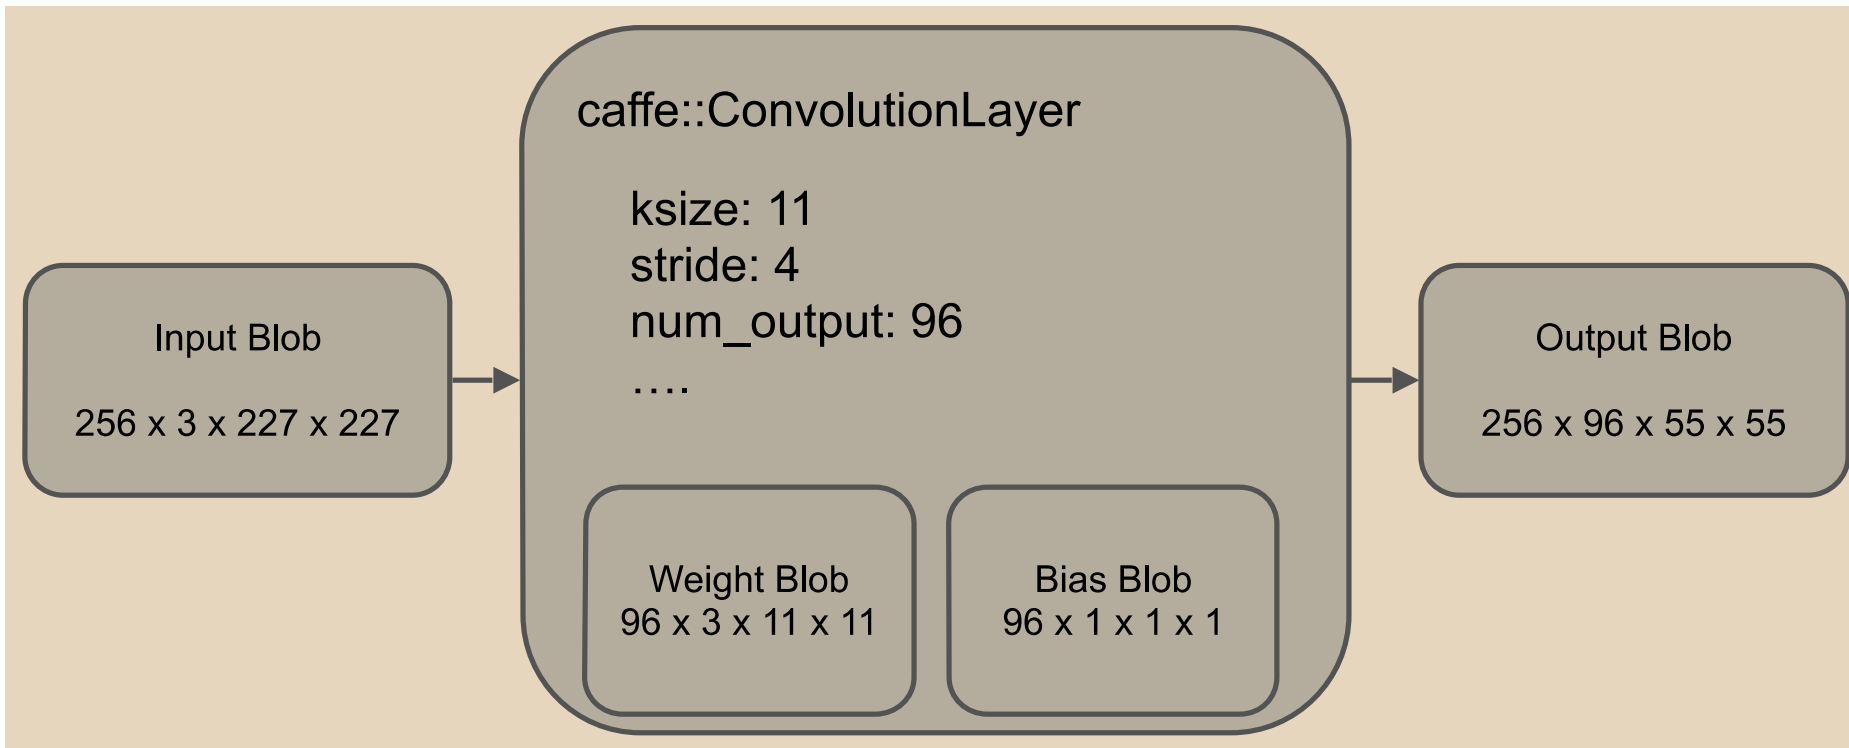


SyncedMem
allocation + communication

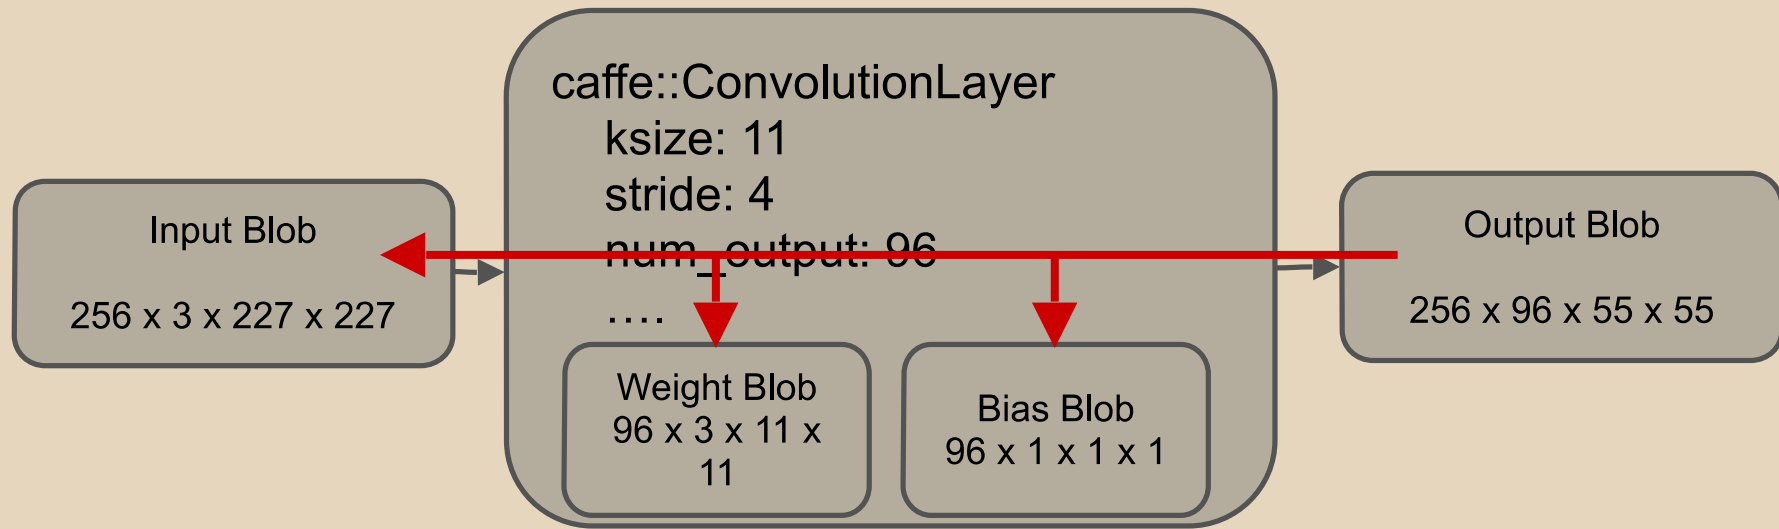


- Earlier, Caffe only supported 4-D blobs and 2-D convolutions ($N \times C \times H \times W$)
- Since October'15, it supports
 - n-D blobs and
 - (n-2)-D convolutions






- Forward: given input, computes the output. →
- Backward: given the gradient w.r.t. the output, compute the gradient w.r.t. the input and its internal parameters. →
- Setup: how to initialize the layer.





Can also have..



In-place
updates

Example: ReLU or PReLU



(PReLU) He *et al.* ICCV'15

How much memory would a PReLU require?

- It does an in-place update, so say requires B for blob
- Say it requires P for parameters (could be per-channel, or just a single scalar)
- Does it need any more?
 - Yes! Need to keep the original input around for computing the derivative for parameters => $+B$
- Q: Can parameterized layers do in-place updates?

GPU/CPU Switch with Blob

- Use synchronized memory
- Mutable/non-mutable determines whether to copy. Use of `mutable_*` may lead to data copy
- Rule of thumb:
Use ~~mutable~~`_`{cpu|gpu}_data whenever possible

Layers

```
src/caffe/layers/absval_layer.cpp
src/caffe/layers/accuracy_layer.cpp
src/caffe/layers/argmax_layer.cpp
src/caffe/layers/base_data_layer.cpp
src/caffe/layers/bnll_layer.cpp
src/caffe/layers/concat_layer.cpp
src/caffe/layers/contrastive_loss_layer.cpp
src/caffe/layers/conv_layer.cpp
src/caffe/layers/cudnn_conv_layer.cpp
src/caffe/layers/cudnn_pooling_layer.cpp
src/caffe/layers/cudnn_relu_layer.cpp
src/caffe/layers/cudnn_sigmoid_layer.cpp
src/caffe/layers/cudnn_softmax_layer.cpp
src/caffe/layers/cudnn_tanh_layer.cpp
src/caffe/layers/data_layer.cpp
src/caffe/layers/dropout_layer.cpp
[abhi@aurora]~/research/codes/caffe-latest :> ls src/caffe/layers/*.cu
src/caffe/layers/absval_layer.cu
src/caffe/layers/base_data_layer.cu
src/caffe/layers/bnll_layer.cu
src/caffe/layers/concat_layer.cu
src/caffe/layers/contrastive_loss_layer.cu
src/caffe/layers/conv_layer.cu
src/caffe/layers/cudnn_conv_layer.cu
src/caffe/layers/cudnn_pooling_layer.cu
src/caffe/layers/cudnn_relu_layer.cu
src/caffe/layers/cudnn_sigmoid_layer.cu
src/caffe/layers/cudnn_softmax_layer.cu
src/caffe/layers/cudnn_tanh_layer.cu
src/caffe/layers/data_layer.cu
src/caffe/layers/dropout_layer.cu
src/caffe/layers/eltwise_layer.cu
src/caffe/layers/euclidean_loss_layer.cu
src/caffe/layers/flatten_layer.cu
src/caffe/layers/hdf5_data_layer.cu
src/caffe/layers/hdf5_output_layer.cu
src/caffe/layers/im2col_layer.cu
src/caffe/layers/inner_product_layer.cu
src/caffe/layers/lrn_layer.cu
src/caffe/layers/mvn_layer.cu
src/caffe/layers/multinomial_logistic_loss_layer.cu
src/caffe/layers/mvnc_layer.cu
src/caffe/layers/neuron_layer.cu
src/caffe/layers/pooling_layer.cu
src/caffe/layers/power_layer.cu
src/caffe/layers/relu_layer.cu
src/caffe/layers/sigmoid_cross_entropy_loss_layer.cu
src/caffe/layers/sigmoid_layer.cu
src/caffe/layers/silence_layer.cu
src/caffe/layers/slice_layer.cu
src/caffe/layers/softmax_layer.cu
src/caffe/layers/softmax_loss_layer.cu
src/caffe/layers/split_layer.cu
src/caffe/layers/tanh_layer.cu
src/caffe/layers/threshold_layer.cu
src/caffe/layers/threshold_loss_layer.cu
src/caffe/layers/window_data_layer.cu
```

More about Layers

- Data layers
- Vision layers
- Common layers
- Activation/Neuron layers
- Loss layers

Data Layers

- Data enters through data layers -- they lie at the bottom of nets.
- Data can come from efficient databases (*LevelDB* or LMDB), directly from memory, or, when efficiency is not critical, from files on disk in HDF5/.mat or common image formats.
- Common input preprocessing (mean subtraction, scaling, random cropping, and mirroring) is available by specifying `TransformationParameters`.

Data Layers

- Data (Backend: LevelDB, LMDB)
- MemoryData
- HDF5Data
- ImageData
- WindowData
- DummyData
- Write your own! In Python!

Data Layers

Database

- Layer type: Data
- Parameters
 - Required
 - `source`: the name of the directory containing the database
 - `batch_size`: the number of inputs to process at one time
 - Optional
 - `rand_skip`: skip up to this number of inputs at the beginning; useful for asynchronous sgd
 - `backend` [default LEVELDB]: choose whether to use a LEVELDB or LMDB

Data Layers

```
1  name: "LeNet"
2  layer {
3    name: "mnist"
4    type: "Data"
5    top: "data"
6    top: "label"
7    include {
8      phase: TRAIN
9    }
10   transform_param {
11     scale: 0.00390625
12   }
13   data_param {
14     source: "examples/mnist/mnist_train_lmdb"
15     batch_size: 64
16     backend: LMDB
17   }
18 }
```


Data Layers

In-Memory

- Layer type: `MemoryData`
- Parameters
 - Required
 - `batch_size, channels, height, width`: specify the size of input chunks to read from memory

The memory data layer reads data directly from memory, without copying it. In order to use it, one must call `MemoryDataLayer::Reset` (from C++) or `Net.set_input_arrays` (from Python) in order to specify a source of contiguous data (as 4D row major array), which is read one batch-sized chunk at a time.

HDF5 Input

- Layer type: HDF5Data
- Parameters
 - Required
 - `source`: the name of the file to read from
 - `batch_size`

HDF5 Output

- Layer type: HDF5Output
- Parameters
 - Required
 - `file_name`: name of file to write to

The HDF5 output layer performs the opposite function of the other layers in this section: it writes its input blobs to disk.

Images

- Layer type: `ImageData`
- Parameters
 - Required
 - `source`: name of a text file, with each line giving an image filename and label
 - `batch_size`: number of images to batch together
 - Optional
 - `rand_skip`
 - `shuffle` [default false]
 - `new_height, new_width`: if provided, resize all images to this size

Windows

`WindowData`

Dummy

`DummyData` is for development and debugging. See `DummyDataParameter`.

Writing your own data layer in python

- Compile Caffe, uncommenting in Makefile.config

```
# WITH_PYTHON_LAYER := 1
```

- Example: See Fast-RCNN

Prototxt

```
1 name: "CaffeNet"
2 layer {
3   name: 'data'
4   type: 'Python'
5   top: 'data'
6   top: 'rois'
7   top: 'labels'
8   top: 'bbox_targets'
9   top: 'bbox_loss_weights'
10  python_param {
11    module: 'roi_data_layer.layer'
12    layer: 'RoIDataLayer'
13    param_str: "'num_classes': 21"
14  }
15 }
```

Python

```
import caffe

class RoIDataLayer(caffe.Layer):
    """Fast R-CNN data layer used for training."""

    def setup(self, bottom, top):
        """Setup the RoIDataLayer."""
        # ...
        pass

    def forward(self, bottom, top):
        # ...
        pass

    def backward(self, top, propagate_down, bottom):
        """This layer does not propagate gradients."""
        pass

    def reshape(self, bottom, top):
        """Reshaping happens during the call to fwd."""
        pass
```

Transformations

```
layer {  
  name: "data"  
  type: "Data"  
  [...]  
  transform_param {  
    scale: 0.1  
    mean_file_size: mean.binaryproto  
    # for images in particular horizontal mirroring and random cropping  
    # can be done as simple data augmentations.  
    mirror: 1 # 1 = on, 0 = off  
    # crop a `crop_size` x `crop_size` patch:  
    # - at random during training  
    # - from the center during testing  
    crop_size: 227  
  }  
}
```

Note that all layers do not support transformations, like HDF5

More about Layers

- Data layers
- Vision layers
- Common layers
- Activation/Neuron layers
- Loss layers

Vision Layers

- Images as input and produce other images as output.
- Non-trivial height $h > 1$ and width $w > 1$.
- 2D geometry naturally lends itself to certain decisions about how to process the input.
 - Since Oct'15, supports nD convolutions
- In particular, most of the vision layers work by applying a particular operation to some region of the input to produce a corresponding region of the output.
- In contrast, other layers (with few exceptions) ignore the spatial structure of the input, effectively treating it as “one big vector” with dimension “***chw***”.

Convolution Layer

Convolution

- Layer type: Convolution
- CPU implementation: `./src/caffe/layers/convolution_layer.cpp`
- CUDA GPU implementation: `./src/caffe/layers/convolution_layer.cu`
- Parameters (ConvolutionParameter convolution_param)
 - Required
 - `num_output (c_o)`: the number of filters
 - `kernel_size` (OR `kernel_h` and `kernel_w`): specifies height and width of each filter
 - Strongly Recommended
 - `weight_filler` [default type: 'constant' value: 0]
 - Optional
 - `bias_term` [default true]: specifies whether to learn and apply a set of additive biases to the filter outputs
 - `pad` (OR `pad_h` and `pad_w`) [default 0]: specifies the number of pixels to (implicitly) add to each side of the input
 - `stride` (OR `stride_h` and `stride_w`) [default 1]: specifies the intervals at which to apply the filters to the input
 - `group (g)` [default 1]: If $g > 1$, we restrict the connectivity of each filter to a subset of the input. Specifically, the input and output channels are separated into g groups, and the i th output group channels will be only connected to the i th input group channels.

Input

- $n * c_i * h_i * w_i$

Output

- $n * c_o * h_o * w_o$, where $h_o = (h_i + 2 * pad_h - kernel_h) / stride_h + 1$ and w_o likewise.

```
layer {
  name: "conv1"
  type: "Convolution"
  bottom: "data"
  top: "conv1"
  # learning rate and decay multipliers for the filters
  param { lr_mult: 1 decay_mult: 1 }
  # learning rate and decay multipliers for the biases
  param { lr_mult: 2 decay_mult: 0 }
  convolution_param {
    num_output: 96      # learn 96 filters
    kernel_size: 11     # each filter is 11x11
    stride: 4           # step 4 pixels between each filter application
    weight_filler {
      type: "gaussian" # initialize the filters from a Gaussian
      std: 0.01        # distribution with stdev 0.01 (default mean: 0)
    }
    bias_filler {
      type: "constant" # initialize the biases to zero (0)
      value: 0
    }
  }
}
}
```

Pooling Layer

- LayerType: POOLING
- CPU implementation: `./src/caffe/layers/pooling_layer.cpp`
- CUDA GPU implementation: `./src/caffe/layers/pooling_layer.cu`
- Parameters (`PoolingParameter pooling_param`)
 - Required
 - `kernel_size` (or `kernel_h` and `kernel_w`): specifies height and width of each filter
 - Optional
 - `pool` [default MAX]: the pooling method. Currently MAX, AVE, or STOCHASTIC
 - `pad` (or `pad_h` and `pad_w`) [default 0]: specifies the number of pixels to (implicitly) add to each side of the input
 - `stride` (or `stride_h` and `stride_w`) [default 1]: specifies the intervals at which to apply the filters to the input

```
layers {  
  name: "pool1"  
  type: POOLING  
  bottom: "conv1"  
  top: "pool1"  
  pooling_param {  
    pool: MAX  
    kernel_size: 3 # pool over a 3x3 region  
    stride: 2      # step two pixels (in the bottom blob) between pooling regions  
  }  
}
```

Vision Layers

- Convolution
- Pooling
- Local Response Normalization (LRN)
- Im2col -- helper

More about Layers

- Data layers
- Vision layers
- Common layers
- Activation/Neuron layers
- Loss layers

Common Layers

- INNER_PRODUCT $W^T x + b$ (fully connected)
- SPLIT
- FLATTEN
- CONCAT
- SLICE
- ELTWISE (element wise operations)
- ARGMAX
- SOFTMAX
- MVN (mean-variance normalization)

```
layer {
  name: "fc8"
  type: "InnerProduct"
  # learning rate and decay multipliers for the weights
  param { lr_mult: 1 decay_mult: 1 }
  # learning rate and decay multipliers for the biases
  param { lr_mult: 2 decay_mult: 0 }
  inner_product_param {
    num_output: 1000
    weight_filler {
      type: "gaussian"
      std: 0.01
    }
    bias_filler {
      type: "constant"
      value: 0
    }
  }
}
bottom: "fc7"
top: "fc8"
}
```

More about Layers

- Data layers
- Vision layers
- Common layers
- Activation/Neuron layers
- Loss layers

Activation/Neuron layers

- One Input Blob
- One Output Blob
 - Both same size
- Or a single blob – in-place updates

Activation/Neuron layers

- ReLU / PReLU
- Sigmoid
- Tanh
- Absval
- Power
- BNLL (binomial normal log likelihood)

ReLU / Rectified-Linear and Leaky-ReLU

- Layer type: ReLU
- CPU implementation: `./src/caffe/layers/relu_layer.cpp`
- CUDA GPU implementation: `./src/caffe/layers/relu_layer.cu`
- Parameters (ReLUParameter relu_param)
 - Optional
 - `negative_slope` [default 0]: specifies whether to leak the negative part by multiplying it with the slope value rather than setting it to 0.
- Sample (as seen in `./models/bvlc_reference_caffenet/train_val.prototxt`)

```
layer {  
  name: "relu1"  
  type: "ReLU"  
  bottom: "conv1"  
  top: "conv1"  
}
```


Given an input value x , The ReLU layer computes the output as x if $x > 0$ and $\text{negative_slope} * x$ if $x \leq 0$. When the negative slope parameter is not set, it is equivalent to the standard ReLU function of taking $\max(x, 0)$. It also supports in-place computation, meaning that the bottom and the top blob could be the same to preserve memory consumption.

More about Layers

- Data layers
- Vision layers
- Common layers
- Activation/Neuron layers
- Loss layers



Loss

What kind of model is this?



Loss

What kind of model is this?



Who knows! Need a **loss function**.

Classification

SOFTMAX_LOSS

HINGE_LOSS

Linear Regression

EUCLIDEAN_LOSS

Attributes / Multiclassification

SIGMOID_CROSS_ENTROPY_LOSS

Others...

New Task

NEW_LOSS

Loss

- **Loss function** determines the learning task.
- Given data D , a Net typically minimizes:

$$L(W) = \frac{1}{|D|} \sum_i^{|D|} f_W(X^{(i)}) + \lambda r(W)$$

Data term: error
averaged over instances

Regularization
term: penalize
large weights
to improve
generalization

Loss

- The data error term $f_W(X^{(i)})$ is computed by `Net::Forward`
- Loss is computed as the output of Layers
- Pick the loss to suit the task – many different losses for different needs

Loss Layers

- `SOFTMAX_LOSS`
- `HINGE_LOSS`
- `EUCLIDEAN_LOSS`
- `SIGMOID_CROSS_ENTROPY_LOSS`
- `INFOGAIN_LOSS`

- `ACCURACY`
- `TOPK`

Loss Layers

- SOFTMAX_LOSS
- HINGE_LOSS
- EUCLIDEAN_LOSS
- SIGMOID_..._LOSS
- INFOGAIN_LOSS
- ACCURACY
- TOPK
- **NEW_LOSS**

Classification

Linear Regression

Attributes /

Multiclassification

Other losses

Not a loss

Softmax Loss Layer

- Multinomial logistic regression: used for predicting a single class of K mutually exclusive classes

```
layers {  
  name: "loss"  
  type: "SoftmaxWithLoss"  
  bottom: "pred"  
  bottom: "label"  
  top: "loss"  
}
```

$$\hat{p}_{nk} = \exp(x_{nk}) / [\sum_{k'} \exp(x_{nk'})]$$

$$E = -\frac{1}{N} \sum_{n=1}^N \log(\hat{p}_{n, l_n})$$

Sigmoid Cross-Entropy Loss

- Binary logistic regression: used for predicting K independent probability values in [0, 1]

```
layers {  
  name: "loss"  
  type: "SigmoidCrossEntropyLoss"  
  bottom: "pred"  
  bottom: "label"  
  top: "loss"  
}
```

$$y = (1 + \exp(-x))^{-1}$$

$$E = \frac{1}{n} \sum_{n=1}^N [p_n \log \hat{p}_n + (1 - p_n) \log(1 - \hat{p}_n)]$$

Euclidean Loss

- A loss for regressing to real-valued labels $[-\text{inf}, \text{inf}]$

```
layers {  
  name: "loss"  
  type: "EuclideanLoss"  
  bottom: "pred"  
  bottom: "label"  
  top: "loss"  
}
```

$$E = \frac{1}{2N} \sum_{n=1}^N \|\hat{y}_n - y_n\|_2^2$$

Multiple loss layers

- Your network can contain as many loss functions as you want, as long as it is a DAG!
- Reconstruction and Classification:

$$E = \frac{1}{2N} \sum_{n=1}^N \|\hat{y}_n - y_n\|_2^2 + \frac{1}{N} \sum_{n=1}^N \log(\hat{p}_{n,l_n}),$$

```
layers {  
  name: "recon-loss"  
  type: "EuclideanLoss"  
  bottom: "reconstructions"  
  bottom: "data"  
  top: "recon-loss"  
}
```

```
layers {  
  name: "class-loss"  
  type: "SoftmaxWithLoss"  
  bottom: "class-preds"  
  bottom: "class-labels"  
  top: "class-loss"  
}
```

Multiple loss layers

“*Loss” layers have a default loss weight of 1

```
layers {  
  name: "loss"  
  type: "SoftmaxWithLoss"  
  bottom: "pred"  
  bottom: "label"  
  top: "loss"  
}
```

==

```
layers {  
  name: "loss"  
  type: "SoftmaxWithLoss"  
  bottom: "pred"  
  bottom: "label"  
  top: "loss"  
  loss_weight: 1.0  
}
```

Multiple loss layers

- Give each loss its own weight
- E.g. give higher priority to classification error
- Or, to balance the values of different loss functions

$$E = \frac{1}{2N} \sum_{n=1}^N \|\hat{y}_n - y_n\|_2^2 + 100 * \frac{1}{N} \sum_{n=1}^N \log(\hat{p}_{n,l_n}),$$

```
layers {  
  name: "recon-loss"  
  type: "EuclideanLoss"  
  bottom: "reconstructions"  
  bottom: "data"  
  top: "recon-loss"  
}
```

```
layers {  
  name: "class-loss"  
  type: "SoftmaxWithLoss"  
  bottom: "class-preds"  
  bottom: "class-labels"  
  top: "class-loss"  
  loss_weight: 100.0
```


Any layer can produce a loss!

- Just add `loss_weight: 1.0` to have a layer's output be incorporated into the loss

$$E = || \text{pred} - \text{label} ||^2 / (2N)$$

```
layers {  
  name: "loss"  
  type: "EuclideanLoss"  
  bottom: "pred"  
  bottom: "label"  
  top: "euclidean_loss"  
  loss_weight: 1.0  
}
```

$$\text{diff} = \text{pred} - \text{label}$$

```
layers {  
  name: "diff"  
  type: "Eltwise"  
  bottom: "pred"  
  bottom: "label"  
  top: "diff"  
  eltwise_param {  
    op: SUM  
    coeff: 1  
    coeff: -1  
  }  
}
```

$$E = || \text{diff} ||^2 / (2N)$$

```
layers {  
  name: "loss"  
  type: "Power"  
  bottom: "diff"  
  top: "euclidean_loss"  
  power_param {  
    power: 2  
  }  
  # = 1/(2N)  
  loss_weight: 0.0078125  
}
```

Layers

- Data layers
- Vision layers
- Common layers
- Activation/Neuron layers
- Loss layers

Initialization

- Gaussian [most commonly used]
 - Xavier
 - Constant [default]
-
- Goal: keep the variance roughly fixed

Solving: Training a Net

Optimization like model definition is configuration.

`train_net: "lenet_train.prototxt"`

`base_lr: 0.01`

`momentum: 0.9`

`weight_decay: 0.0005`

`max_iter: 10000`

`snapshot_prefix: "lenet_snapshot"`

All you need to run things
on the GPU.



`> caffe train -solver lenet_solver.prototxt -gpu 0`

Stochastic Gradient Descent (SGD) + momentum •

Adaptive Gradient (ADAGRAD) • Nesterov's Accelerated Gradient (NAG)

```
# The train/test net protocol buffer definition
net: "logreg_train.prototxt"
# test_iter specifies how many forward passes the test should carry out.
# In the case of MNIST, we have test batch size 100 and 100 test iterations,
# covering the full 10,000 testing images.
test_iter: 100
# Carry out testing every 500 training iterations.
test_interval: 500
# The base learning rate, momentum and the weight decay of the network.
base_lr: 0.01
momentum: 0.9
weight_decay: 0.0005
# The learning rate policy
lr_policy: "inv"
gamma: 0.0001
power: 0.75
# Display every 100 iterations
display: 100
# The maximum number of iterations
max_iter: 10000
# snapshot intermediate results
snapshot: 5000
snapshot_prefix: "examples/mnist/lenet"
# solver mode: CPU or GPU
solver_mode: GPU
```

Solver

- **Solver** optimizes the network weights W to minimize the loss $L(W)$ over the data D

$$L(W) = \frac{1}{|D|} \sum_i^{|D|} f_W \left(X^{(i)} \right) + \lambda r(W)$$

- Coordinates forward / backward, weight updates, and scoring.

Solver

- Computes parameter update ΔW formed from
 - The stochastic error gradient ∇f_W
 - The regularization gradient $\nabla r(W)$
 - Particulars to each solving method

$$L(W) \approx \frac{1}{N} \sum_i^N f_W (X^{(i)}) + \lambda r(W)$$

SGD Solver

- Stochastic gradient descent, with momentum
- `solver_type: SGD`

$$V_{t+1} = \mu V_t - \alpha \nabla L(W_t)$$

$$W_{t+1} = W_t + V_{t+1}$$

SGD Solver

- “AlexNet” [1] training strategy:
 - Use momentum 0.9
 - Initialize learning rate at 0.01
 - Periodically drop learning rate by a factor of 10
- Just a few lines of Caffe solver specification:

```
base_lr: 0.01
lr_policy: "step"
gamma: 0.1
stepsize: 100000
max_iter: 350000
momentum: 0.9
```

NAG Solver

- Nesterov's accelerated gradient [1]
- `solver_type: NESTEROV`
- Proven to have optimal convergence rate $\mathcal{O}(1/t^2)$ for convex problems

$$V_{t+1} = \mu V_t - \alpha \nabla L(W_t + \mu V_t)$$

$$W_{t+1} = W_t + V_{t+1}$$

AdaGrad Solver

- Adaptive gradient (Duchi et al. [1])
- `solver_type: ADAGRAD`
- Attempts to automatically scale gradients based on historical gradients

$$(W_{t+1})_i = (W_t)_i - \alpha \frac{(\nabla L(W_t))_i}{\sqrt{\sum_{t'=1}^t (\nabla L(W_{t'}))_i^2}}$$

Solver Showdown: MNIST Autoencoder

AdaGrad

```
I0901 13:36:30.007884 24952 solver.cpp:232] Iteration 65000, loss = 64.1627
I0901 13:36:30.007922 24952 solver.cpp:251] Iteration 65000, Testing net (#0) # train set
I0901 13:36:33.019305 24952 solver.cpp:289] Test loss: 63.217
I0901 13:36:33.019356 24952 solver.cpp:302] Test net output #0: cross_entropy_loss = 63.217 (* 1 = 63.217 loss)
I0901 13:36:33.019773 24952 solver.cpp:302] Test net output #1: l2_error = 2.40951
```

SGD

```
I0901 13:35:20.426187 20072 solver.cpp:232] Iteration 65000, loss = 61.5498
I0901 13:35:20.426218 20072 solver.cpp:251] Iteration 65000, Testing net (#0) # train set
I0901 13:35:22.780092 20072 solver.cpp:289] Test loss: 60.8301
I0901 13:35:22.780138 20072 solver.cpp:302] Test net output #0: cross_entropy_loss = 60.8301 (* 1 = 60.8301 loss)
I0901 13:35:22.780146 20072 solver.cpp:302] Test net output #1: l2_error = 2.02321
```

Nesterov

```
I0901 13:36:52.466069 22488 solver.cpp:232] Iteration 65000, loss = 59.9389
I0901 13:36:52.466099 22488 solver.cpp:251] Iteration 65000, Testing net (#0) # train set
I0901 13:36:55.068370 22488 solver.cpp:289] Test loss: 59.3663
I0901 13:36:55.068410 22488 solver.cpp:302] Test net output #0: cross_entropy_loss = 59.3663 (* 1 = 59.3663 loss)
I0901 13:36:55.068418 22488 solver.cpp:302] Test net output #1: l2_error = 1.79998
```

Weight sharing

- Parameters can be shared and reused across Layers throughout the Net
- Applications:
 - Convolution at multiple scales / pyramids
 - Recurrent Neural Networks (RNNs)
 - Siamese nets for distance learning

Weight sharing

- Just give the parameter blobs explicit names using the `param` field
- Layers specifying the same `param` name will share that parameter, accumulating gradients accordingly

```
layers: {
  name: 'innerproduct1'
  type: "InnerProduct"
  inner_product_param {
    num_output: 10
    bias_term: false
    weight_filler {
      type: 'gaussian'
      std: 10
    }
  }
  param: 'sharedweights'
  bottom: 'data'
  top: 'innerproduct1'
}
layers: {
  name: 'innerproduct2'
  type: "InnerProduct"
  inner_product_param {
    num_output: 10
    bias_term: false
  }
  param: 'sharedweights'
  bottom: 'data'
  top: 'innerproduct2'
}
```

Interfaces

- Command Line
- Python
- Matlab

CMD

\$> Caffe --params

```
# train LeNet
caffe train -solver examples/mnist/lenet_solver.prototxt
# train on GPU 2
caffe train -solver examples/mnist/lenet_solver.prototxt -gpu 2
# resume training from the half-way point snapshot
caffe train -solver examples/mnist/lenet_solver.prototxt -snapshot
examples/mnist/lenet_iter_5000.solverstate
```


CMD

```
# fine-tune CaffeNet model weights for pascal
caffe train \
-solver examples/finetune_pascal_detection/pascal_finetune_solver.prototxt -
weights models/bvlc_reference_caffenet/bvlc_reference_caffenet.caffemodel

# score the learned LeNet model on the validation set as defined in the model
architecture lenet_train_test.prototxt
caffe test -model examples/mnist/lenet_train_test.prototxt -weights
examples/mnist/lenet_iter_10000 -gpu 0 -iterations 100
```

CMD

```
# (These example calls require you complete the LeNet / MNIST example first.)  
# time LeNet training on CPU for 10 iterations  
caffe time -model examples/mnist/lenet_train_test.prototxt -iterations 10  
# time a model architecture with the given weights on the first GPU for 10  
iterations  
# time LeNet training on GPU for the default 50 iterations  
caffe time -model examples/mnist/lenet_train_test.prototxt -gpu 0  
  
# query the first device  
caffe device_query -gpu 0
```

Python

```
$> make pycaffe
```

```
python> import caffe
```

`caffe.Net`: is the central interface for loading, configuring, and running models.

`caffe.Classssifier` & `caffe.Detector` for convenience

`caffe.SGDSolver` exposes the solving interface.

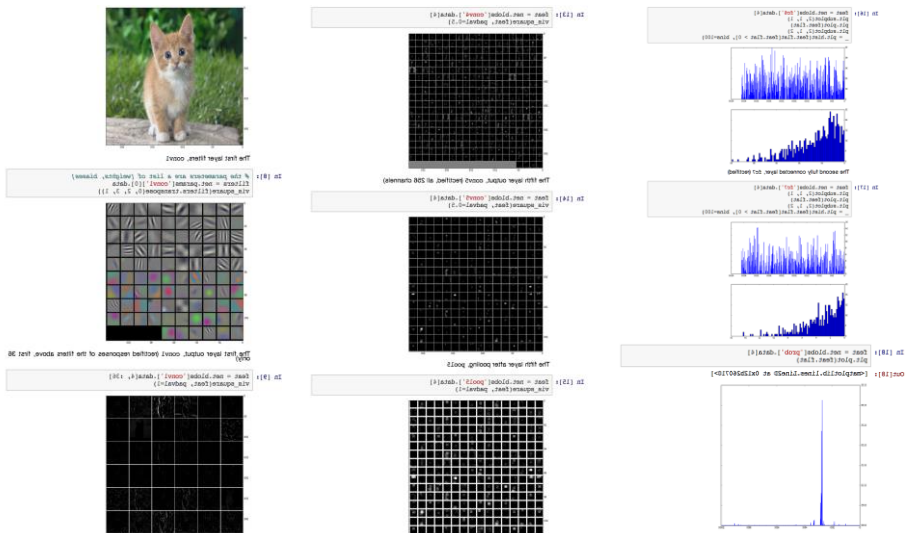
`caffe.io` handles I/O with preprocessing and protocol buffers.

`caffe.draw` visualizes network architectures.

`Caffe blobs` are exposed as numpy ndarrays for ease-of-use and efficiency**

Python

GOTO: [IPython Filter Visualization Notebook](#)



MATLAB

```
522 static handler_registry handlers[] = {
523     // Public API functions
524     { "get_solver",      get_solver      },
525     { "solver_get_attr", solver_get_attr },
526     { "solver_get_iter", solver_get_iter },
527     { "solver_restore",  solver_restore  },
528     { "solver_solve",    solver_solve    },
529     { "solver_step",     solver_step     },
530     { "get_net",         get_net         },
531     { "net_get_attr",    net_get_attr    },
532     { "net_forward",     net_forward     },
533     { "net_backward",    net_backward    },
534     { "net_copy_from",   net_copy_from   },
535     { "net_reshape",     net_reshape     },
536     { "net_save",        net_save        },
537     { "layer_get_attr",  layer_get_attr },
538     { "layer_get_type",  layer_get_type  },
539     { "blob_get_shape",  blob_get_shape },
540     { "blob_reshape",    blob_reshape    },
541     { "blob_get_data",   blob_get_data   },
542     { "blob_set_data",   blob_set_data   },
543     { "blob_get_diff",   blob_get_diff   },
544     { "blob_set_diff",   blob_set_diff   },
545     { "set_mode_cpu",    set_mode_cpu    },
546     { "set_mode_gpu",    set_mode_gpu    },
547     { "set_device",      set_device      },
548     { "get_init_key",    get_init_key    },
549     { "reset",           reset           },
550     { "read_mean",       read_mean       },
551     { "write_mean",      write_mean      },
552     { "version",         version         },
553     // The end.
554     { "END",            NULL            },
555 };
```

RECENT MODELS


- Network-in-Network (NIN)
- GoogLeNet
- VGG

Questions?


THAT'S ALL! THANKS!

Network-in-Network

- filter with a nonlinear composition instead of a linear filter
- 1x1 convolution + nonlinearity
- reduce dimensionality, deepen the representation




Linear Filter
CONV

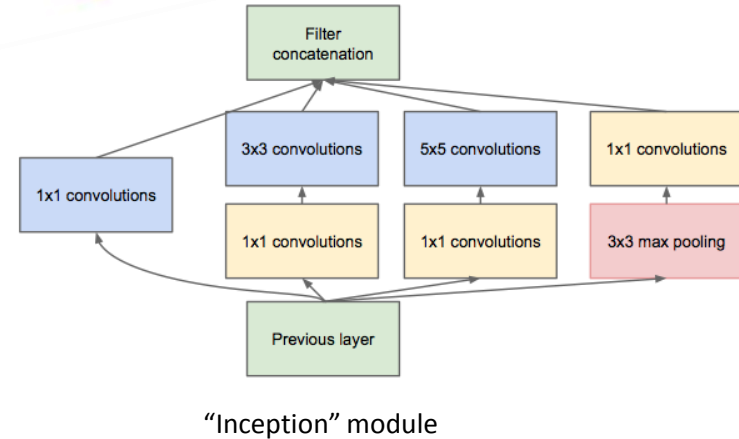


NIN / MLP filter
1x1 CONV

GoogLeNet



- composition of multi-scale dimension-reduced “Inception” modules
- 1x1 conv for dimensionality reduction
- concatenation across filter scales
- multiple losses for training to depth



VGG

- 3x3 convolution all the way down...
- fine-tuned progression of deeper models
- 16 and 19 parameter layer variations in the model zoo

ConvNet Configuration					
A	A-LRN	B	C	D	E
11 weight layers	11 weight layers	13 weight layers	16 weight layers	16 weight layers	19 weight layers
input (224 × 224 RGB image)					
conv3-64	conv3-64 LRN	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64
maxpool					
conv3-128	conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128
maxpool					
conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256 conv1-256	conv3-256 conv3-256 conv3-256	conv3-256 conv3-256 conv3-256 conv3-256
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
FC-4096					
FC-4096					
FC-1000					
soft-max					

Table 2: **Number of parameters** (in millions).

Network	A, A-LRN	B	C	D	E
Number of parameters	133	133	134	138	144

Blob Data Management

```
// Assuming that data are on the CPU initially, and we have a blob.  
const Dtype* foo;  
Dtype* bar;  
foo = blob.gpu_data(); // data copied cpu->gpu.  
foo = blob.cpu_data(); // no data copied since both have up-to-date contents.  
bar = blob.mutable_gpu_data(); // no data copied.  
// ... some operations ...  
bar = blob.mutable_gpu_data(); // no data copied when we are still on GPU.  
foo = blob.cpu_data(); // data copied gpu->cpu, since the gpu side has modified the data  
foo = blob.gpu_data(); // no data copied since both have up-to-date contents  
bar = blob.mutable_cpu_data(); // still no data copied.  
bar = blob.mutable_gpu_data(); // data copied cpu->gpu.  
bar = blob.mutable_cpu_data(); // data copied gpu->cpu.
```