

# Introduction to Machine Learning

## CMU-10701

### Deep Learning

Barnabás Póczos & Aarti Singh



**MACHINE LEARNING** DEPARTMENT



# Credits

Many of the pictures, results, and other materials are taken from:

Ruslan Salakhutdinov

Joshua Bengio

Geoffrey Hinton

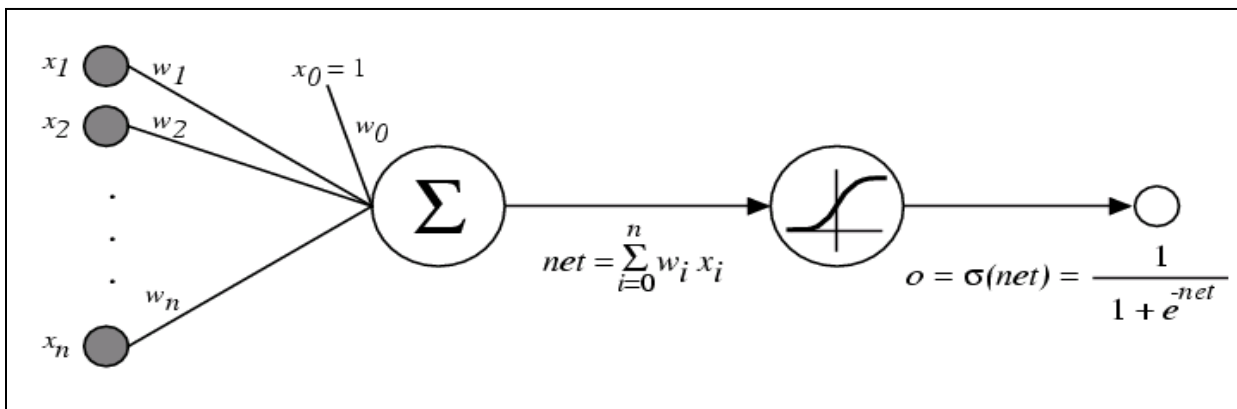

Yann LeCun

# Contents

- ❑ Definition and Motivation
- ❑ History of Deep architectures
- ❑ Deep architectures
  - ❑ Convolutional networks
  - ❑ Deep Belief networks
- ❑ Applications

# Deep architectures

**Defintion:** Deep architectures are composed of *multiple levels* of non-linear operations, such as neural nets with many hidden layers.



# Goal of Deep architectures

**Goal:** Deep learning methods aim at

- learning *feature hierarchies*
- where features from higher levels of the hierarchy are formed by lower level features.

edges, local shapes, object parts

Low level representation

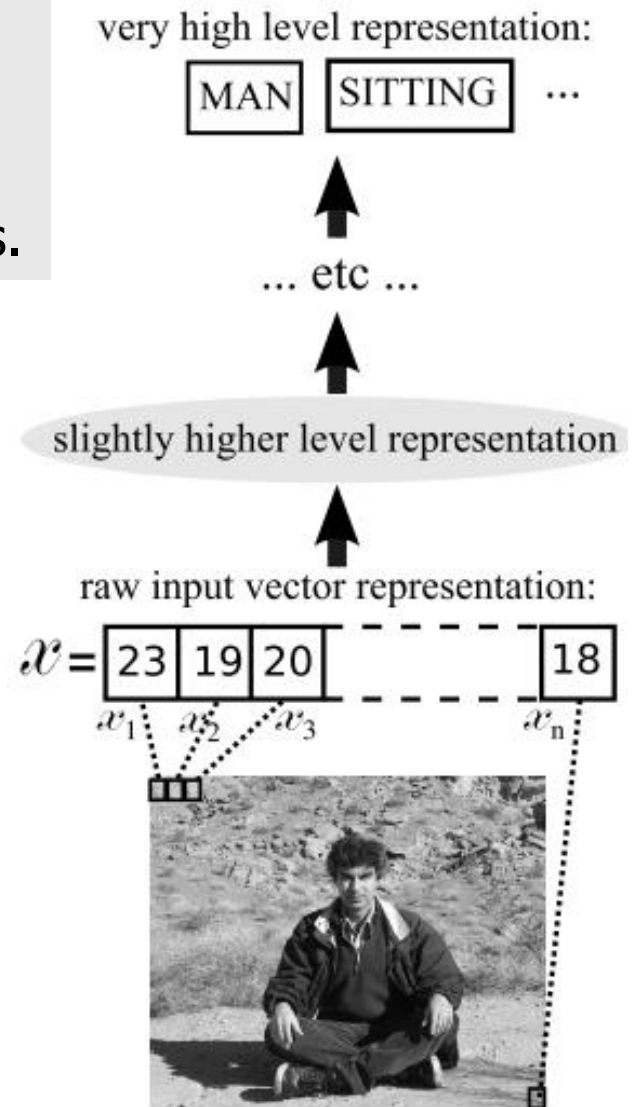



Figure is from Yoshua Bengio

# Neurobiological Motivation

- ❑ Most current learning algorithms are shallow architectures (1-3 levels)  
(SVM, kNN, MoG, KDE, Parzen Kernel regression, PCA, Perceptron,...)

$$\text{SVM: } \hat{f}(\mathbf{x}) = \text{sign}\left(\sum_{i=1}^n \alpha_i y_i k(\mathbf{x}_i, \mathbf{x})\right)$$



- ❑ The mammal brain is organized in a deep architecture (Serre, Kreiman, Kouh, Cadieu, Knoblich, & Poggio, 2007)  
(E.g. visual system has 5 to 10 levels)

# Deep Learning History

- ❑ **Inspired** by the architectural depth of the brain, researchers wanted for decades to train deep multi-layer neural networks.
- ❑ **No successful** attempts were reported before 2006 ...
  - Researchers reported positive experimental results with typically two or three levels (i.e. one or two hidden layers), but training deeper networks consistently yielded poorer results.
- ❑ **Exception:** convolutional neural networks, LeCun 1998
- ❑ **SVM:** Vapnik and his co-workers developed the Support Vector Machine (1993). It is a shallow architecture.
- ❑ **Digression:** In the 1990's, many researchers abandoned neural networks with multiple adaptive hidden layers because SVMs worked better, and there was no successful attempts to train deep networks.
- ❑ **Breakthrough in 2006**

# Breakthrough

## **Deep Belief Networks (DBN)**

Hinton, G. E, Osindero, S., and Teh, Y. W. (2006).  
A fast learning algorithm for deep belief nets.  
Neural Computation, 18:1527-1554.

## **Autoencoders**

Bengio, Y., Lamblin, P., Popovici, P., Larochelle, H. (2007).  
Greedy Layer-Wise Training of Deep Networks,  
Advances in Neural Information Processing Systems 19



# Theoretical Advantages of Deep Architectures

- ❑ Some functions cannot be efficiently represented (in terms of number of tunable elements) by architectures that are too shallow.
- ❑ Deep architectures might be able to represent some functions otherwise not efficiently representable.


- ❑ **More formally:**

Functions that can be compactly represented by a depth  $k$  architecture might require an exponential number of computational elements to be represented by a depth  $k - 1$  architecture

- ❑ The consequences are
  - **Computational:** We don't need exponentially many elements in the layers
  - **Statistical:** poor generalization may be expected when using an insufficiently deep architecture for representing some functions.

# Theoretical Advantages of Deep Architectures

## The Polynomial circuit:



# Deep Convolutional Networks

# Deep Convolutional Networks

- ❑ Deep supervised neural networks are generally too difficult to train.
- ❑ **One notable exception:** convolutional neural networks (CNN)
- ❑ Convolutional nets were inspired by the visual system's structure
- ❑ They typically have five, six or seven layers, a number of layers which makes fully-connected neural networks almost impossible to train properly when initialized randomly.

# Deep Convolutional Networks


Compared to standard feedforward neural networks with similarly-sized layers,

- CNNs have much fewer connections and parameters
- and so they are easier to train,
- while their theoretically-best performance is likely to be only slightly worse.

## LeNet 5


Y. LeCun, L. Bottou, Y. Bengio and P. Haffner: **Gradient-Based Learning Applied to Document Recognition**, *Proceedings of the IEEE*, 86(11):2278-2324, November **1998**

# LeNet 5, LeCun 1998



- Input: 32x32 pixel image. Largest character is 20x20 (All important info should be in the center of the receptive field of the highest level feature detectors)
- Cx: Convolutional layer
- Sx: Subsample layer
- Fx: Fully connected layer
- Black and White pixel values are normalized:  
E.g. White = -0.1, Black = 1.175 (Mean of pixels = 0, Std of pixels = 1)

# LeNet 5, Layer C1



C1: Convolutional layer with 6 feature maps of size 28x28.  $C1_k$  ( $k=1\dots6$ )

Each unit of C1 has a 5x5 receptive field in the input layer.

- Topological structure
- Sparse connections
- Shared weights


$(5*5+1)*6=156$  parameters to learn

Connections:  $28*28*(5*5+1)*6=122304$

If it was fully connected we had  $(32*32+1)*(28*28)*6$  parameters



# LeNet 5, Layer S2




S2: Subsampling layer with 6 feature maps of size 14x14  
2x2 nonoverlapping receptive fields in C1

Layer S2:  $6 \times 2 = 12$  trainable parameters.

Connections:  $14 \times 14 \times (2 \times 2 + 1) \times 6 = 5880$



# LeNet 5, Layer C3



- C3: Convolutional layer with 16 feature maps of size 10x10
- Each unit in C3 is connected to several! 5x5 receptive fields at identical locations in S2

Layer C3:

1516 trainable parameters.


Connections: 151600

|   | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 | 15 |
|---|---|---|---|---|---|---|---|---|---|---|----|----|----|----|----|----|
| 0 | X |   |   |   | X | X | X |   |   | X | X  | X  | X  |    | X  | X  |
| 1 | X | X |   |   |   | X | X | X |   |   | X  | X  | X  | X  |    | X  |
| 2 | X | X | X |   |   |   | X | X | X |   |    | X  |    | X  | X  | X  |
| 3 |   | X | X | X |   |   | X | X | X | X |    |    | X  |    | X  | X  |
| 4 |   |   | X | X | X |   |   | X | X | X | X  |    | X  | X  |    | X  |
| 5 |   |   |   | X | X | X |   |   | X | X | X  | X  |    | X  | X  | X  |

TABLE I

EACH COLUMN INDICATES WHICH FEATURE MAP IN S2 ARE COMBINED BY THE UNITS IN A PARTICULAR FEATURE MAP OF C3.

# LeNet 5, Layer S4




- S4: Subsampling layer with 16 feature maps of size 5x5
- Each unit in S4 is connected to the corresponding 2x2 receptive field at C3

Layer S4:  $16 \times 2 = 32$  trainable parameters.

Connections:  $5 \times 5 \times (2 \times 2 + 1) \times 16 = 2000$


# LeNet 5, Layer C5



- C5: Convolutional layer with 120 feature maps of size 1x1
- Each unit in C5 is connected to all 16 5x5 receptive fields in S4

Layer C5:  $120 \times (16 \times 25 + 1) = 48120$  trainable parameters and connections  
(Fully connected)

# LeNet 5, Layer C5



Layer F6: 84 fully connected units.  $84 \times (120 + 1) = 10164$  trainable parameters and connections.

Output layer: 10RBF (One for each digit)

84=7x12, stylized image


**Weight update:** Backpropagation

# MINIST Dataset



60,000 original datasets

Test error: 0.95%



540,000 artificial distortions

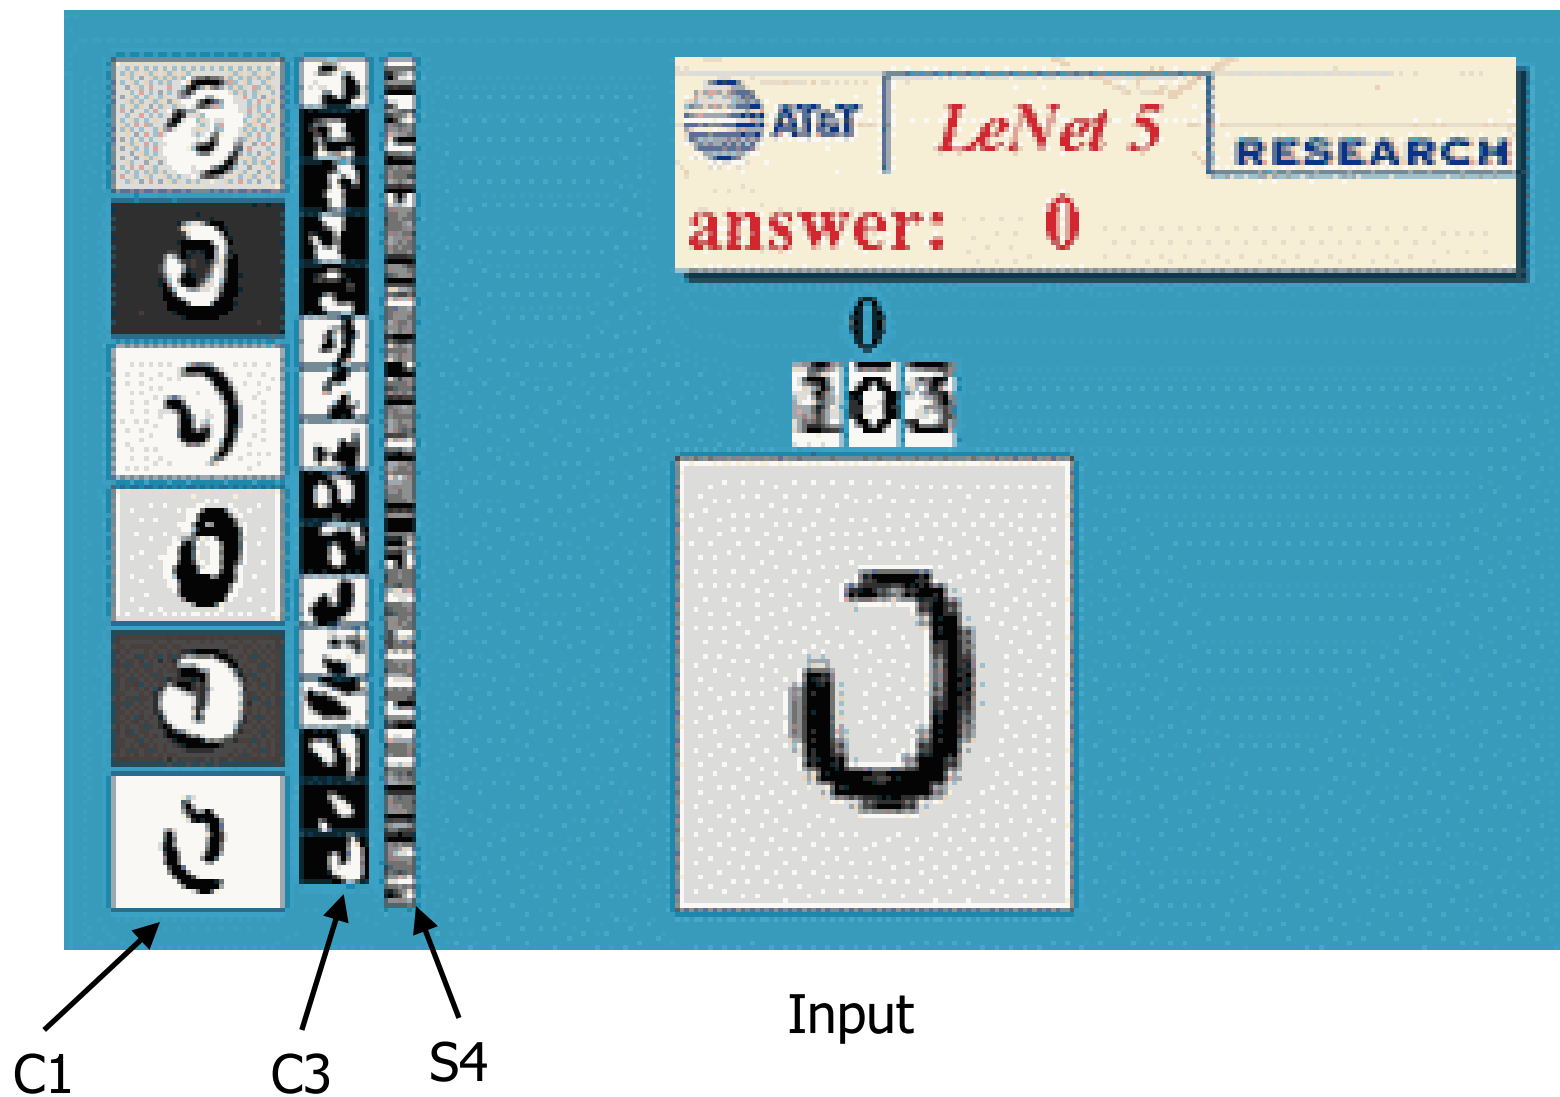
+ 60,000 original

Test error: 0.8%

# Misclassified examples



# LeNet 5 in Action




# LeNet 5, Shift invariance








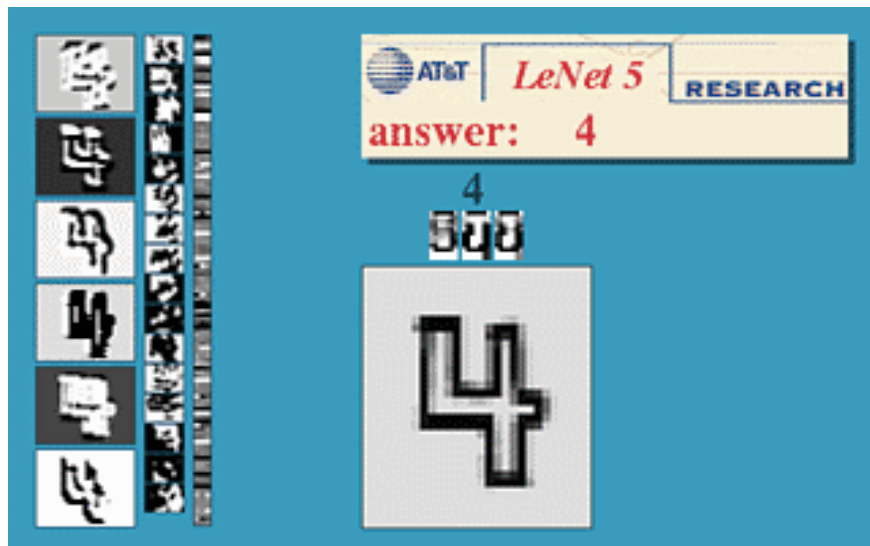
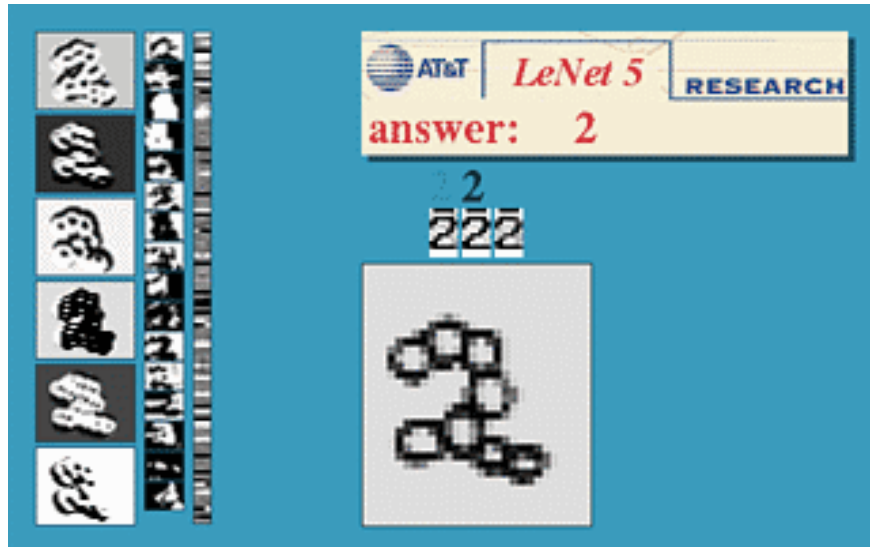
# LeNet 5, Rotation invariance



# LeNet 5, Noise resistance



# LeNet 5, Unusual Patterns



# ImageNet Classification with Deep Convolutional Neural Networks

Alex Krizhevsky, Ilya Sutskever, Geoffrey Hinton,  
Advances in Neural Information Processing Systems 2012

# ImageNet

- ❑ 15M images
- ❑ 22K categories
- ❑ Images collected from Web
- ❑ Human labelers (Amazon's Mechanical Turk crowd-sourcing)
- ❑ ImageNet Large Scale Visual Recognition Challenge (ILSVRC-2010)
  - 1K categories
  - 1.2M training images (~1000 per category)
  - 50,000 validation images
  - 150,000 testing images
- ❑ RGB images
- ❑ Variable-resolution, but this architecture scales them to 256x256 size

# ImageNet

## Classification goals:

- ❑ Make 1 guess about the label (Top-1 error)
- ❑ make 5 guesses about the label (Top-5 error)




# The Architecture

Typical nonlinearities:  $f(x) = \tanh(x)$

$$f(x) = (1 + e^{-x})^{-1}$$


Here, however, Rectified Linear Units (ReLU) are used:  $f(x) = \max(0, x)$

**Empirical observation:** Deep convolutional neural networks with ReLUs train several times faster than their equivalents with tanh units



A four-layer convolutional neural network with ReLUs (solid line) reaches a 25% training error rate on CIFAR-10 six times faster than an equivalent network with tanh neurons (dashed line)

# The Architecture



**The first convolutional layer** filters the  $224 \times 224 \times 3$  input image with 96 kernels of size  $11 \times 11 \times 3$  with a stride of 4 pixels (this is the distance between the receptive field centers of neighboring neurons in the kernel map.  $224/4=56$

**The pooling layer:** form of non-linear down-sampling. Max-pooling partitions the input image into a set of rectangles and, for each such sub-region, outputs the maximum value



# The Architecture

- Trained with stochastic gradient descent
  - on two NVIDIA GTX 580 3GB GPUs
  - for about a week
- 
- ❑ 650,000 neurons
  - ❑ 60,000,000 parameters
  - ❑ 630,000,000 connections
  - ❑ 5 convolutional layer, 3 fully connected layer
  - ❑ Final feature layer: 4096-dimensional

# Data Augmentation

The easiest and most common method to **reduce overfitting** on image data is to artificially **enlarge the dataset** using label-preserving transformations.

We employ two distinct forms of data augmentation:

- image translation
- horizontal reflections
- changing RGB intensities

# Dropout

- ❑ We know that combining different models can be very useful (Mixture of experts, majority voting, boosting, etc)
- ❑ Training many different models, however, is very time consuming.

## **The solution:**

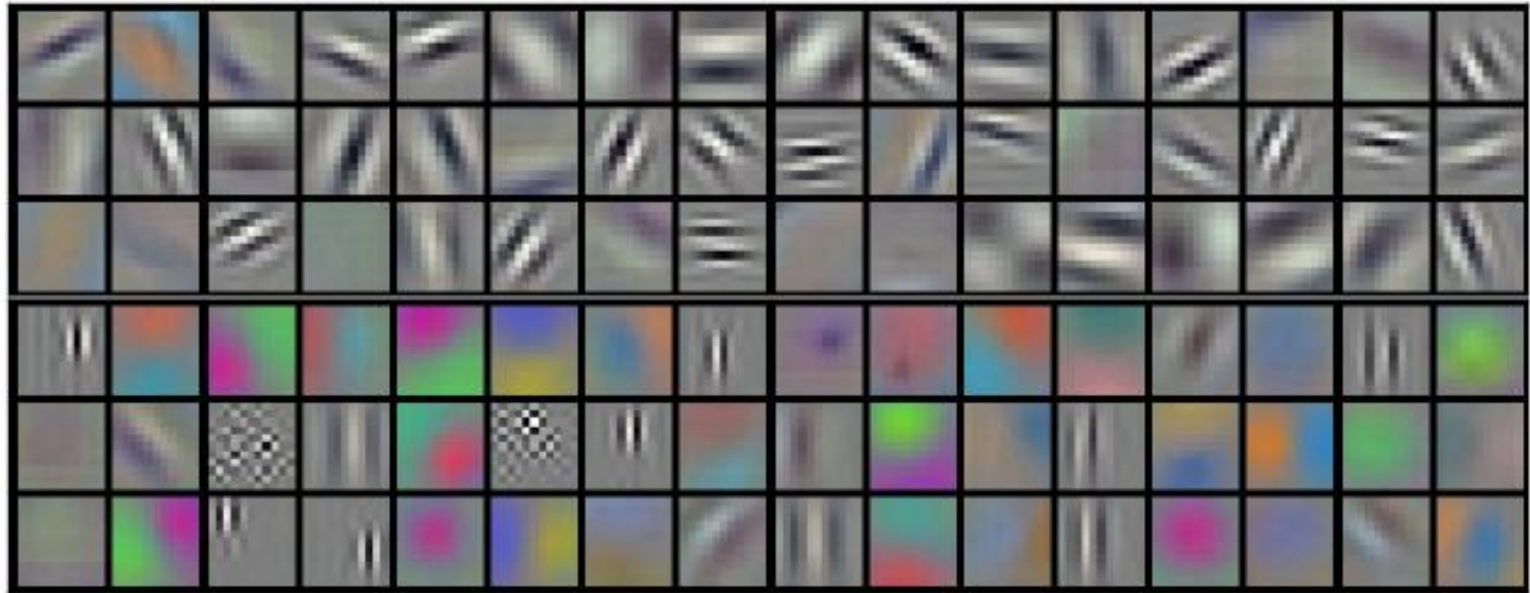
*Dropout*: set the output of each hidden neuron to zero w.p. 0.5.

# Dropout

**Dropout:** set the output of each hidden neuron to zero w.p. 0.5.

- The neurons which are “dropped out” in this way do not contribute to the forward pass and do not participate in backpropagation.
- So every time an input is presented, the neural network samples a different architecture, but all these architectures share weights.
- This technique reduces complex co-adaptations of neurons, since a neuron cannot rely on the presence of particular other neurons.
- It is, therefore, forced to learn more robust features that are useful in conjunction with many different random subsets of the other neurons.
- Without dropout, our network exhibits substantial overfitting.
- Dropout roughly doubles the number of iterations required to converge.

# The first convolutional layer



96 convolutional kernels of size  $11 \times 11 \times 3$  learned by the first convolutional layer on the  $224 \times 224 \times 3$  input images.

The top 48 kernels were learned on GPU1 while the bottom 48 kernels were learned on GPU2

Looks like Gabor wavelets, ICA filters...

# Results

## **Results on the test data:**

top-1 error rate: 37.5%

top-5 error rate: 17.0%

## **ILSVRC-2012 competition:**

15.3% accuracy

2<sup>nd</sup> best team: 26.2% accuracy

# Results



**mite**

**container ship**

**motor scooter**

**leopard**

|  |  |   |  |
|--|--|---|--|
|  |  |   |  |
| mite<br>black widow<br>cockroach<br>tick<br>starfish | container ship<br>lifeboat<br>amphibian<br>fireboat<br>drilling platform | motor scooter<br>go-kart<br>moped<br>bumper car<br>golfcart | leopard<br>jaguar<br>cheetah<br>snow leopard<br>Egyptian cat |



**grille**

**mushroom**

**cherry**

**Madagascar cat**

|   |   |   |  |
|---|---|---|--|
|   |   |   |  |
| convertible<br>grille<br>pickup<br>beach wagon<br>fire engine | agaric<br>mushroom<br>jelly fungus<br>gill fungus<br>dead-man's-fingers | dalmatian<br>grape<br>elderberry<br>ffordshire bullterrier<br>currant | squirrel monkey<br>spider monkey<br>titi<br>indri<br>howler monkey |



# Results: Image similarity



Test column

six training images that produce feature vectors in the last hidden layer with the smallest Euclidean distance from the feature vector for the test image.



# Deep Belief Networks

# What is wrong with back propagation?

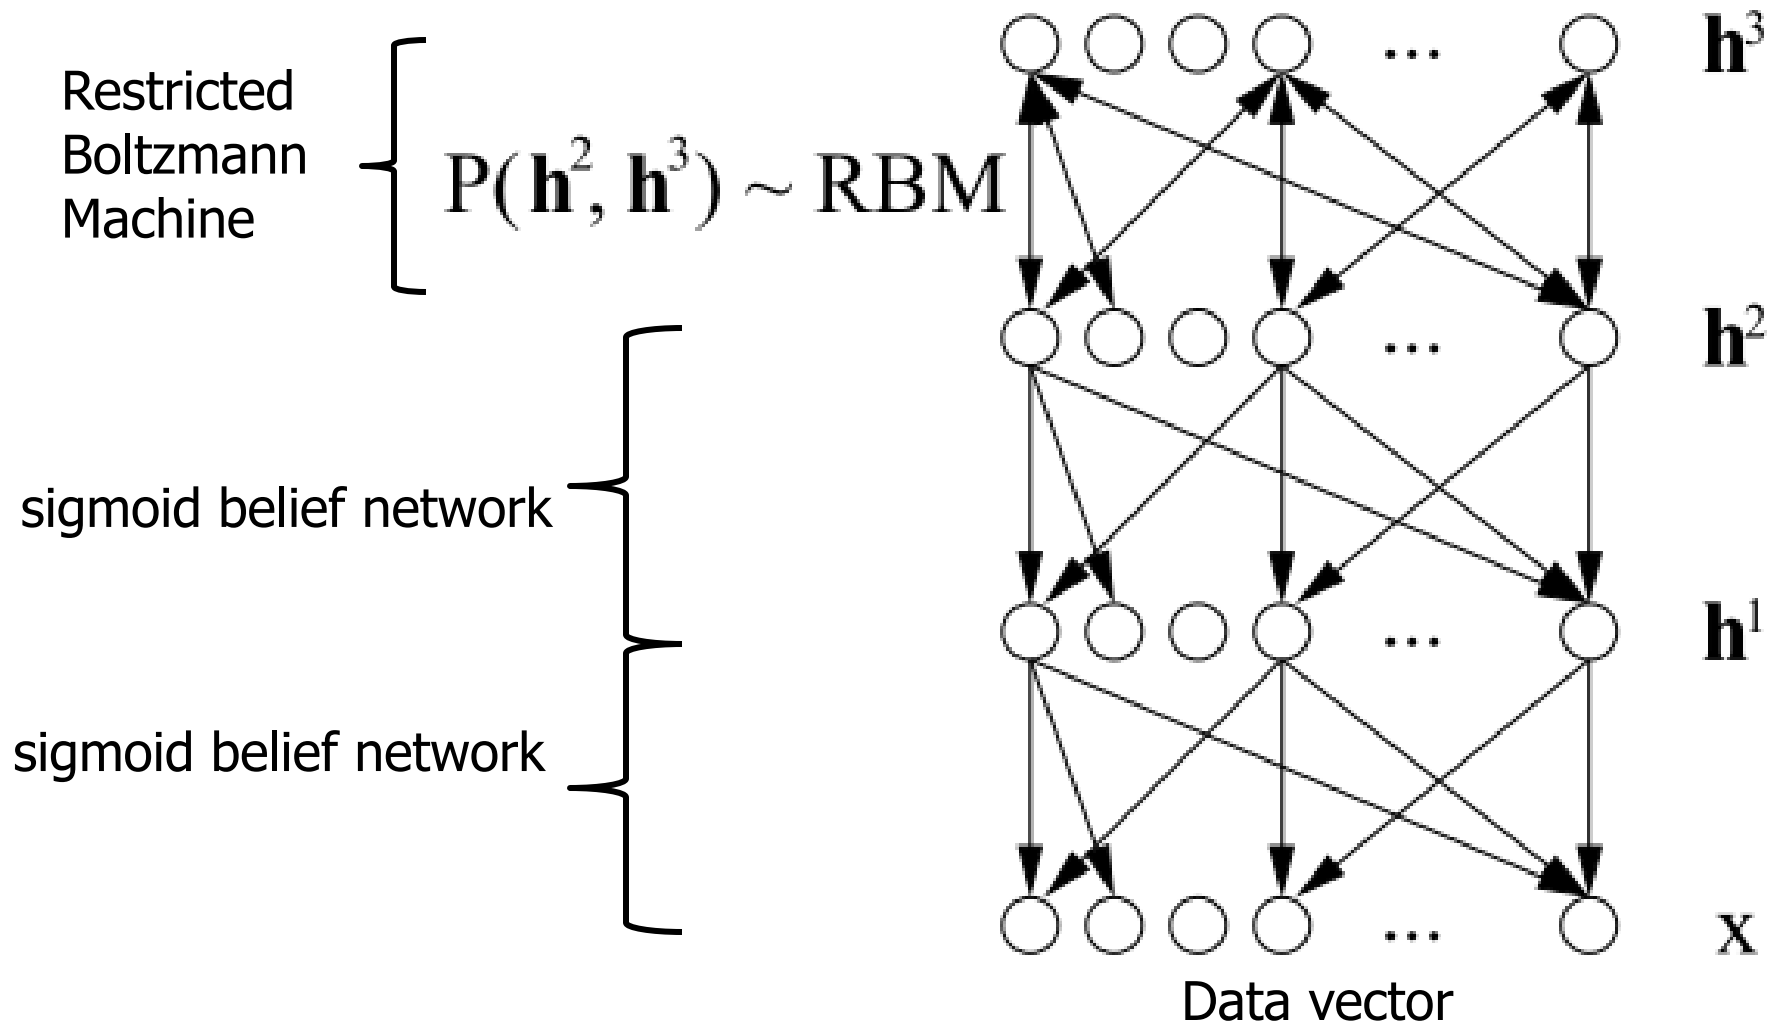
- ❑ It requires labeled training data.
  - Almost all data is unlabeled.
- ❑ The learning time does not scale well.
  - It is very slow in networks with multiple hidden layers.
- ❑ It can get stuck in poor local optima.
  - Usually in deep nets they are far from optimal.
- ❑ MLP is not a generative model, it only focuses on  $P(Y|X)$ .  
We would like a generative approach that could learn  $P(X)$  as well.
- ❑ **Solution:** *Deep Belief Networks*, a generative graphical model

# Deep Belief Network

## **Deep Belief Networks (DBN's)**

- are probabilistic generative models
- contain many layers of hidden variables
- each layer captures high-order correlations between the activities of hidden features in the layer below
- the top two layers of the DBN form an undirected bipartite graph called Restricted Boltzmann Machine
- the lower layers forming a directed sigmoid belief network

# Deep Belief Network



# Deep Belief Network


$$P(\mathbf{h}^l, \mathbf{h}^{l-1}) \propto \exp(\mathbf{b}^T \mathbf{h}^{l-1} + \mathbf{c}^T \mathbf{h}^l + \mathbf{h}^{lT} W \mathbf{h}^{l-1})$$

$$P(\mathbf{h}^2, \mathbf{h}^3) \sim \text{RBM}$$

$$\sigma(x) = (1 + e^{-x})^{-1}$$

$$P(h_i^k = 1 | \mathbf{h}^{k+1}) = \sigma(b_i^{k+1} + \sum_j W_{i,j}^{k+1} h_j^{k+1})$$

$$P(x_i = 1 | \mathbf{h}^1) = \sigma(b_i^1 + \sum_j W_{i,j}^1 h_j^1)$$



**Joint likelihood:**

$$P(\mathbf{x}, \mathbf{h}^1, \dots, \mathbf{h}^l) = P(\mathbf{h}^l, \mathbf{h}^{l-1}) \left( \prod_{k=1}^{l-2} P(\mathbf{h}^k | \mathbf{h}^{k+1}) \right) P(\mathbf{x} | \mathbf{h}^1)$$

# Boltzmann Machines

# Boltzmann Machines

**Boltzmann machine:** a network of symmetrically coupled stochastic binary units  $\{0,1\}$

## Parameters:

$$\theta = \{\mathbf{W}, \mathbf{L}, \mathbf{J}\}$$

$\mathbf{W}$ : visible-to-hidden

$\mathbf{L}$ : visible-to-visible,  $\text{diag}(\mathbf{L})=0$


$\mathbf{J}$ : hidden-to-hidden,  $\text{diag}(\mathbf{J})=0$

Hidden layer

$$\mathbf{h} \in \{0, 1\}^p$$

Visible layer

$$\mathbf{v} \in \{0, 1\}^d$$



**Energy of the Boltzmann machine:**

$$E(\mathbf{v}, \mathbf{h}|\theta) = -\frac{1}{2}\mathbf{v}^T \mathbf{L} \mathbf{v} - \frac{1}{2}\mathbf{h}^T \mathbf{J} \mathbf{h} - \mathbf{v}^T \mathbf{W} \mathbf{h}$$

# Boltzmann Machines

**Energy of the Boltzmann machine:**

$$E(\mathbf{v}, \mathbf{h}|\theta) = -\frac{1}{2}\mathbf{v}^T \mathbf{L} \mathbf{v} - \frac{1}{2}\mathbf{h}^T \mathbf{J} \mathbf{h} - \mathbf{v}^T \mathbf{W} \mathbf{h}$$

**Generative model:**

Joint likelihood:  $P(\mathbf{v}, \mathbf{h}|\theta) \propto \exp(-E(\mathbf{v}, \mathbf{h}; \theta))$

**Probability of a visible vector  $\mathbf{v}$ :**

$$P(\mathbf{v}|\theta) = \frac{\sum_{\mathbf{h}} \exp(-E(\mathbf{v}, \mathbf{h}; \theta))}{Z(\theta)}$$

**Exponentially large set**

$$Z(\theta) = \sum_{\mathbf{v}} \sum_{\mathbf{h}} \exp(-E(\mathbf{v}, \mathbf{h}; \theta))$$



# Restricted Boltzmann Machines

No hidden-to-hidden and no visible-to-visible connections.

$W$ : visible-to-hidden

$L = 0$ : visible-to-visible

$J = 0$ : hidden-to-hidden

**Energy of RBM:**


$$E(\mathbf{v}, \mathbf{h}|\theta) = -\mathbf{v}^T W \mathbf{h} - \mathbf{b}^T \mathbf{v} - \mathbf{a}^T \mathbf{h}$$

**Joint likelihood:**


$$P(\mathbf{v}, \mathbf{h}|\theta) = \frac{1}{Z(\theta)} \exp(-E(\mathbf{v}, \mathbf{h}; \theta))$$

Hidden layer

Visible layer



# Restricted Boltzmann Machines



Top layer: vector of stochastic binary hidden units **h**

Bottom layer: a vector of stochastic binary visible variables **v**.

# Training RBM

Due to the special bipartite structure of RBM's, the hidden units can be explicitly marginalized out:

$$P(\mathbf{v}; \theta) = \frac{1}{\mathcal{Z}(\theta)} \sum_{\mathbf{h}} \exp(-E(\mathbf{v}, \mathbf{h}; \theta)).$$

$$\begin{aligned} P(\mathbf{v}; \theta) &= \frac{1}{\mathcal{Z}(\theta)} \sum_{\mathbf{h}} \exp(\mathbf{v}^\top W \mathbf{h} + \mathbf{b}^\top \mathbf{v} + \mathbf{a}^\top \mathbf{h}) \\ &= \frac{1}{\mathcal{Z}(\theta)} \exp(\mathbf{b}^\top \mathbf{v}) \prod_{j=1}^F \sum_{h_j \in \{0,1\}} \exp\left(a_j h_j + \sum_{i=1}^D W_{ij} v_i h_j\right) \\ &= \frac{1}{\mathcal{Z}(\theta)} \exp(\mathbf{b}^\top \mathbf{v}) \prod_{j=1}^F \left(1 + \exp\left(a_j + \sum_{i=1}^D W_{ij} v_i\right)\right). \end{aligned}$$

# Training RBM

$$P(\mathbf{v}; \theta) = \frac{1}{Z(\theta)} \exp(\mathbf{b}^\top \mathbf{v}) \prod_{j=1}^F \left( 1 + \exp \left( a_j + \sum_{i=1}^D W_{ij} v_i \right) \right)$$

**Gradient descent:**

$$\frac{\partial \log P(\mathbf{v}; \theta)}{\partial W} = E_{P_{\text{data}}}[\mathbf{v}\mathbf{h}^\top] - E_{P_{\text{Model}}}[\mathbf{v}\mathbf{h}^\top],$$

$$\frac{\partial \log P(\mathbf{v}; \theta)}{\partial \mathbf{a}} = E_{P_{\text{data}}}[\mathbf{h}] - E_{P_{\text{Model}}}[\mathbf{h}],$$

$$\frac{\partial \log P(\mathbf{v}; \theta)}{\partial \mathbf{b}} = E_{P_{\text{data}}}[\mathbf{v}] - E_{P_{\text{Model}}}[\mathbf{v}].$$

The exact calculations are intractable because the expectation operator in  $E_{P_{\text{Model}}}$  takes exponential time in  $\min(D, F)$

Efficient Gibbs sampling based approximation exists (Contrastive divergence)

# Inference in RBM

Inference is simple in RBM:

$$P(\mathbf{h}|\mathbf{v}; \theta) = \prod_j p(h_j|\mathbf{v}), \quad P(\mathbf{v}|\mathbf{h}; \theta) = \prod_i p(v_i|\mathbf{h}),$$

$$p(h_j = 1|\mathbf{v}) = g \left( \sum_i W_{ij} v_i + a_j \right),$$

$$p(v_i = 1|\mathbf{h}) = g \left( \sum_j W_{ij} h_j + b_i \right),$$

where  $g(x) = 1/(1 + \exp(-x))$  is the logistic function.

# Training Deep Belief Networks

# Training Deep Belief Networks

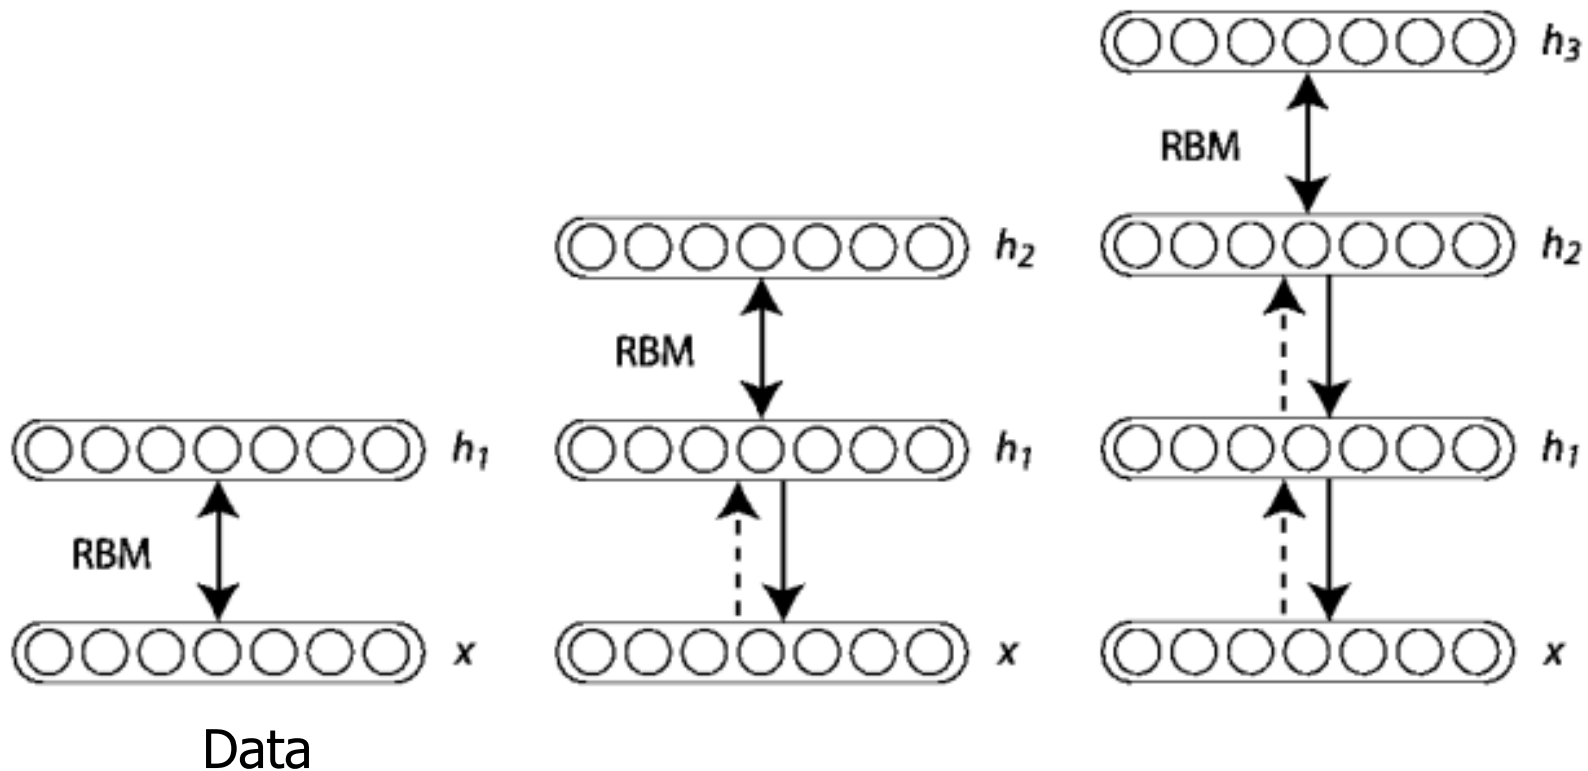
## **Greedy layer-wise unsupervised learning:**

Much better results could be achieved when pre-training each layer with an unsupervised learning algorithm, one layer after the other, starting with the first layer (that directly takes in the observed  $x$  as input).

- The initial experiments used the RBM generative model for each layer.
- Later variants: auto-encoders for training each layer (Bengio et al., 2007; Ranzato et al., 2007; Vincent et al., 2008)
- After having initialized a number of layers, the whole neural network can be fine-tuned with respect to a supervised training criterion as usual

# Training Deep Belief Networks

The unsupervised greedy layer-wise training serves as **initialization**, **replacing** the traditional **random initialization** of multi-layer networks.






# Training Deep Belief Networks


---

**Algorithm 1** Recursive Greedy Learning Procedure for the DBN.

---

- 1: Fit parameters  $W^1$  of the 1<sup>st</sup> layer RBM to data.
  - 2: Freeze the parameter vector  $W^1$  and use samples  $\mathbf{h}^1$  from  $Q(\mathbf{h}^1|\mathbf{v}) = P(\mathbf{h}^1|\mathbf{v}, W^1)$  as the data for training the next layer of binary features with an RBM.
  - 3: Freeze the parameters  $W^2$  that define the 2<sup>nd</sup> layer of features and use the samples  $\mathbf{h}^2$  from  $Q(\mathbf{h}^2|\mathbf{h}^1) = P(\mathbf{h}^2|\mathbf{h}^1, W^2)$  as the data for training the 3<sup>rd</sup> layer of binary features.
  - 4: Proceed recursively for the next layers.
- 






- Deep architecture trained online with 10 million examples of digit images, either with pre-training (triangles) or without (circles).
- The first 2.5 million examples are used for unsupervised pre-training.
- One can see that without pre-training, training converges to a poorer apparent local minimum: unsupervised pre-training helps to find a better minimum of the online error.

# Results

# Deep Boltzmann Machines Results



# Deep Boltzmann Machines Results

2-layer BM



3-layer BM




Training Samples




# Deep Boltzmann Machines Results

## Deep Boltzmann Machine




# Deep Boltzmann Machines Results

Training Samples



Generated Samples



Thanks for your Attention! 😊