

Lecture 6: Particle Filtering — Sequential Importance Resampling and Rao-Blackwellized Particle Filtering

Simo Särkkä


Department of Biomedical Engineering and Computational Science
Aalto University

February 23, 2012

Contents

- 1 Principle of Particle Filter
- 2 Monte Carlo Integration and Importance Sampling
- 3 Sequential Importance Sampling and Resampling
- 4 Rao-Blackwellized Particle Filter
- 5 Particle Filter Properties
- 6 Summary and Demonstration

Particle Filtering: Principle




- Animation: Kalman vs. Particle Filtering:
 - ▶ [Kalman filter animation](#)
 - ▶ [Particle filter animation](#)
- The idea is to form a **weighted particle presentation** $(\mathbf{x}^{(i)}, w^{(i)})$ of the posterior distribution:

$$p(\mathbf{x}) \approx \sum_i w^{(i)} \delta(\mathbf{x} - \mathbf{x}^{(i)}).$$

- Approximates Bayesian optimal filtering equations with **importance sampling**.
- Particle filtering = **Sequential importance sampling**, with additional **resampling** step.

Monte Carlo Integration




- In **Bayesian inference** we often want to compute posterior expectations of the form

$$\mathbb{E}[\mathbf{g}(\mathbf{x}) \mid \mathbf{y}_{1:T}] = \int \mathbf{g}(\mathbf{x}) p(\mathbf{x} \mid \mathbf{y}_{1:T}) d\mathbf{x}$$

- **Monte Carlo**: draw N independent random samples from $\mathbf{x}^{(i)} \sim p(\mathbf{x} \mid \mathbf{y}_{1:T})$ and estimate the expectation as

$$\mathbb{E}[\mathbf{g}(\mathbf{x}) \mid \mathbf{y}_{1:T}] \approx \frac{1}{N} \sum_{i=1}^N \mathbf{g}(\mathbf{x}^{(i)}).$$

Importance Sampling: Basic Version [1/2]



- In practice, we rarely can directly draw samples from the distribution $p(\mathbf{x} \mid \mathbf{y}_{1:T})$.
- In **importance sampling (IS)**, we draw samples from an **importance distribution** $\mathbf{x}^{(i)} \sim \pi(\mathbf{x} \mid \mathbf{y}_{1:T})$ and compute weights $\tilde{w}^{(i)}$ such that

$$\mathbb{E}[\mathbf{g}(\mathbf{x}) \mid \mathbf{y}_{1:T}] \approx \sum_{i=1}^N \tilde{w}^{(i)} \mathbf{g}(\mathbf{x}^{(i)})$$

- **Importance sampling** is based on the identity

$$\begin{aligned} \mathbb{E}[\mathbf{g}(\mathbf{x}) \mid \mathbf{y}_{1:T}] &= \int \mathbf{g}(\mathbf{x}) p(\mathbf{x} \mid \mathbf{y}_{1:T}) d\mathbf{x} \\ &= \int \left[\mathbf{g}(\mathbf{x}) \frac{p(\mathbf{x} \mid \mathbf{y}_{1:T})}{\pi(\mathbf{x} \mid \mathbf{y}_{1:T})} \right] \pi(\mathbf{x} \mid \mathbf{y}_{1:T}) d\mathbf{x} \end{aligned}$$

- Thus we can form a **Monte Carlo approximation** as follows:

$$\mathbb{E}[\mathbf{g}(\mathbf{x}) \mid \mathbf{y}_{1:T}] \approx \frac{1}{N} \sum_{i=1}^N \frac{p(\mathbf{x}^{(i)} \mid \mathbf{y}_{1:T})}{\pi(\mathbf{x}^{(i)} \mid \mathbf{y}_{1:T})} \mathbf{g}(\mathbf{x}^{(i)})$$

- That is, the **importance weights** can be defined as

$$\tilde{w}^{(i)} = \frac{1}{N} \frac{p(\mathbf{x}^{(i)} \mid \mathbf{y}_{1:T})}{\pi(\mathbf{x}^{(i)} \mid \mathbf{y}_{1:T})}$$

Importance Sampling: Weight Normalization

- The problem is that we need to evaluate the **normalization constant** of $p(\mathbf{x}^{(i)} | \mathbf{y}_{1:T})$ – often not possible.
- However, it turns out that we get a valid algorithm if we define **unnormalized importance weights** as

$$w^{*(i)} = \frac{p(\mathbf{y}_{1:T} | \mathbf{x}^{(i)}) p(\mathbf{x}^{(i)})}{\pi(\mathbf{x}^{(i)} | \mathbf{y}_{1:T})}$$

and then **normalize** them:

$$w^{(i)} = \frac{w^{*(i)}}{\sum_j w^{*(j)}}$$

- The (weight-normalized) importance sampling approximation is then

$$\mathbb{E}[\mathbf{g}(\mathbf{x}) | \mathbf{y}_{1:T}] \approx \sum_{i=1}^N w^{(i)} \mathbf{g}(\mathbf{x}^{(i)})$$

Importance Sampling

- Draw N samples from the **importance distribution**:

$$\mathbf{x}^{(i)} \sim \pi(\mathbf{x} \mid \mathbf{y}_{1:T}), \quad i = 1, \dots, N.$$

- Compute the **unnormalized weights** by

$$w^{*(i)} = \frac{p(\mathbf{y}_{1:T} \mid \mathbf{x}^{(i)}) p(\mathbf{x}^{(i)})}{\pi(\mathbf{x}^{(i)} \mid \mathbf{y}_{1:T})},$$

and the **normalized weights** by

$$w^{(i)} = \frac{w^{*(i)}}{\sum_{j=1}^N w^{*(j)}}.$$

Importance Sampling: Properties

- The **approximation** to the **posterior expectation** of $\mathbf{g}(\mathbf{x})$ is then given as

$$\mathbb{E}[\mathbf{g}(\mathbf{x}) \mid \mathbf{y}_{1:T}] \approx \sum_{i=1}^N w^{(i)} \mathbf{g}(\mathbf{x}^{(i)}).$$

- The **posterior probability density approximation** can then be formally written as

$$p(\mathbf{x} \mid \mathbf{y}_{1:T}) \approx \sum_{i=1}^N w^{(i)} \delta(\mathbf{x} - \mathbf{x}^{(i)}),$$

where $\delta(\cdot)$ is the Dirac delta function.

- The **efficiency** depends on the choice of the **importance distribution**.

Sequential Importance Sampling: Idea

- Sequential Importance Sampling (SIS) is concerned with models

$$\mathbf{x}_k \sim p(\mathbf{x}_k \mid \mathbf{x}_{k-1})$$

$$\mathbf{y}_k \sim p(\mathbf{y}_k \mid \mathbf{x}_k)$$

- The SIS algorithm uses a weighted set of *particles* $\{(w_k^{(i)}, \mathbf{x}_k^{(i)}) : i = 1, \dots, N\}$ such that

$$\mathbb{E}[\mathbf{g}(\mathbf{x}_k) \mid \mathbf{y}_{1:k}] \approx \sum_{i=1}^N w_k^{(i)} \mathbf{g}(\mathbf{x}_k^{(i)}).$$

- Or equivalently

$$p(\mathbf{x}_k \mid \mathbf{y}_{1:k}) \approx \sum_{i=1}^N w_k^{(i)} \delta(\mathbf{x}_k - \mathbf{x}_k^{(i)}),$$

where $\delta(\cdot)$ is the Dirac delta function.

- Uses **importance sampling sequentially**.

Sequential Importance Sampling: Derivation [1/2]

- Let's consider the **full posterior** distribution of states $\mathbf{x}_{0:k}$ given the measurements $\mathbf{y}_{1:k}$.
- We get the following **recursion** for the **posterior distribution**:

$$\begin{aligned} p(\mathbf{x}_{0:k} | \mathbf{y}_{1:k}) &\propto p(\mathbf{y}_k | \mathbf{x}_{0:k}, \mathbf{y}_{1:k-1}) p(\mathbf{x}_{0:k} | \mathbf{y}_{1:k-1}) \\ &= p(\mathbf{y}_k | \mathbf{x}_k) p(\mathbf{x}_k | \mathbf{x}_{0:k-1}, \mathbf{y}_{1:k-1}) p(\mathbf{x}_{0:k-1} | \mathbf{y}_{1:k-1}) \\ &= p(\mathbf{y}_k | \mathbf{x}_k) p(\mathbf{x}_k | \mathbf{x}_{k-1}) p(\mathbf{x}_{0:k-1} | \mathbf{y}_{1:k-1}). \end{aligned}$$

- We could now construct an **importance distribution** $\mathbf{x}_{0:k}^{(i)} \sim \pi(\mathbf{x}_{0:k} | \mathbf{y}_{1:k})$ and compute the corresponding (normalized) **importance weights** as

$$w_k^{(i)} \propto \frac{p(\mathbf{y}_k | \mathbf{x}_k^{(i)}) p(\mathbf{x}_k^{(i)} | \mathbf{x}_{k-1}^{(i)}) p(\mathbf{x}_{0:k-1}^{(i)} | \mathbf{y}_{1:k-1})}{\pi(\mathbf{x}_{0:k}^{(i)} | \mathbf{y}_{1:k})}.$$

Sequential Importance Sampling: Derivation [2/2]

- Let's form the **importance distribution recursively** as follows:

$$\pi(\mathbf{x}_{0:k} | \mathbf{y}_{1:k}) = \pi(\mathbf{x}_k | \mathbf{x}_{0:k-1}, \mathbf{y}_{1:k}) \pi(\mathbf{x}_{0:k-1} | \mathbf{y}_{1:k-1})$$

- Expression for the **importance weights** can be written as

$$w_k^{(i)} \propto \frac{p(\mathbf{y}_k | \mathbf{x}_k^{(i)}) p(\mathbf{x}_k^{(i)} | \mathbf{x}_{k-1}^{(i)})}{\pi(\mathbf{x}_k^{(i)} | \mathbf{x}_{0:k-1}, \mathbf{y}_{1:k})} \underbrace{\frac{p(\mathbf{x}_{0:k-1}^{(i)} | \mathbf{y}_{1:k-1})}{\pi(\mathbf{x}_{0:k-1}^{(i)} | \mathbf{y}_{1:k-1})}}_{\propto w_{k-1}^{(i)}}$$

- Thus the weights satisfy the **recursion**

$$w_k^{(i)} \propto \frac{p(\mathbf{y}_k | \mathbf{x}_k^{(i)}) p(\mathbf{x}_k^{(i)} | \mathbf{x}_{k-1}^{(i)})}{\pi(\mathbf{x}_k^{(i)} | \mathbf{x}_{0:k-1}, \mathbf{y}_{1:k})} w_{k-1}^{(i)}$$

Sequential Importance Sampling: Algorithm

Sequential Importance Sampling

- **Initialization:** Draw N samples $\mathbf{x}_0^{(i)}$ from the prior

$$\mathbf{x}_0^{(i)} \sim p(\mathbf{x}_0)$$

and set $w_0^{(i)} = 1/N$.

- **Prediction:** Draw N new samples $\mathbf{x}_k^{(i)}$ from importance distributions

$$\mathbf{x}_k^{(i)} \sim \pi(\mathbf{x}_k | \mathbf{x}_{0:k-1}^{(i)}, \mathbf{y}_{1:k})$$

- **Update:** Calculate new weights according to

$$w_k^{(i)} \propto w_{k-1}^{(i)} \frac{p(\mathbf{y}_k | \mathbf{x}_k^{(i)}) p(\mathbf{x}_k^{(i)} | \mathbf{x}_{k-1}^{(i)})}{\pi(\mathbf{x}_k^{(i)} | \mathbf{x}_{0:k-1}^{(i)}, \mathbf{y}_{1:k})}$$

Sequential Importance Sampling: Degeneracy

- The problem in SIS is that the algorithm is **degenerate**
- It can be shown that the **variance** of the **weights increases at every step**
- It means that we will always converge to **single non-zero weight** $w^{(i)} = 1$ and the rest being zero – not very useful algorithm.
- **Solution:** resampling!

Sequential Importance Resampling: Resampling Step

- Sequential Importance Resampling (SIR) algorithm adds the following resampling step to SIS algorithm:

Resampling

- Interpret each weight $w_k^{(i)}$ as the probability of obtaining the sample index i in the set $\{\mathbf{x}_k^{(i)} \mid i = 1, \dots, N\}$.
 - Draw N samples from that discrete distribution and replace the old sample set with this new one.
 - Set all weights to the constant value $w_k^{(i)} = 1/N$.
- There are many algorithms for implementing this – **stratified resampling** is optimal in terms of variance.

Sequential Importance Resampling: Effective Number of Particles

- A simple way to do resampling is at every step – but every **resampling** operation **increases variance**.
- We can also resample at, say, **every K th step**.
- In **adaptive resampling**, we resample when the effective number of samples is too low (say, $N/10$):

$$n_{\text{eff}} \approx \frac{1}{\sum_{i=1}^N \left(w_k^{(i)}\right)^2},$$

- In theory, **biased**, but in practice works very well and is often used.

Sequential Importance Resampling

- Draw point $\mathbf{x}_k^{(i)}$ from the importance distribution:

$$\mathbf{x}_k^{(i)} \sim \pi(\mathbf{x}_k \mid \mathbf{x}_{0:k-1}^{(i)}, \mathbf{y}_{1:k}), \quad i = 1, \dots, N.$$

- Calculate new weights

$$w_k^{(i)} \propto w_{k-1}^{(i)} \frac{p(\mathbf{y}_k \mid \mathbf{x}_k^{(i)}) p(\mathbf{x}_k^{(i)} \mid \mathbf{x}_{k-1}^{(i)})}{\pi(\mathbf{x}_k^{(i)} \mid \mathbf{x}_{0:k-1}^{(i)}, \mathbf{y}_{1:k})}, \quad i = 1, \dots, N,$$

and normalize them to sum to unity.

- If the effective number of particles is too low, perform resampling.

Sequential Importance Resampling: Bootstrap filter

- In **bootstrap filter** we use the **dynamic model** as the importance distribution

$$\pi(\mathbf{x}_k^{(i)} \mid \mathbf{x}_{0:k-1}, \mathbf{y}_{1:k}) = p(\mathbf{x}_k^{(i)} \mid \mathbf{x}_{k-1}^{(i)})$$

and resample at every step:

Bootstrap Filter

- Draw point $\mathbf{x}_k^{(i)}$ from the dynamic model:

$$\mathbf{x}_k^{(i)} \sim p(\mathbf{x}_k \mid \mathbf{x}_{k-1}^{(i)}), \quad i = 1, \dots, N.$$

- Calculate new weights

$$w_k^{(i)} \propto p(\mathbf{y}_k \mid \mathbf{x}_k^{(i)}), \quad i = 1, \dots, N,$$

and normalize them to sum to unity.

- Perform resampling.

Sequential Importance Resampling: Optimal Importance Distribution

- The **optimal importance** distribution is

$$\pi(\mathbf{x}_k^{(i)} \mid \mathbf{x}_{0:k-1}^{(i)}, \mathbf{y}_{1:k}) = p(\mathbf{x}_k^{(i)} \mid \mathbf{x}_{k-1}^{(i)}, \mathbf{y}_k)$$

- Then the weight update reduces to

$$w_k^{(i)} \propto w_{k-1}^{(i)} p(\mathbf{y}_k \mid \mathbf{x}_{k-1}^{(i)}), \quad i = 1, \dots, N.$$

- The optimal importance distribution can be used, for example, when the state space is finite.

Sequential Importance Resampling: Importance Distribution via Kalman Filtering

- We can also form a **Gaussian approximation** to the optimal importance distribution:

$$p(\mathbf{x}_k^{(i)} \mid \mathbf{x}_{k-1}^{(i)}, \mathbf{y}_k) \approx N(\mathbf{x}_k^{(i)} \mid \tilde{\mathbf{m}}_k^{(i)}, \tilde{\mathbf{P}}_k^{(i)}).$$

by using a single prediction and update steps of a Gaussian filter starting from a singular distribution at $\mathbf{x}_{k-1}^{(i)}$.

- We can also replace above with the result of a Gaussian filter $N(\mathbf{m}_{k-1}^{(i)}, \mathbf{P}_{k-1}^{(i)})$ started from a random initial mean.
- A very common way seems to be to use the previous sample as the mean: $N(\mathbf{x}_{k-1}^{(i)}, \mathbf{P}_{k-1}^{(i)})$.
- A particle filter with UKF proposal has been given name **unscented particle filter (UPF)** – you can invent new PFs easily this way.

Rao-Blackwellized Particle Filter: Idea

- Rao-Blackwellized particle filtering (RBPF) is concerned with **conditionally Gaussian models**:

$$p(\mathbf{x}_k | \mathbf{x}_{k-1}, \boldsymbol{\theta}_{k-1}) = \mathcal{N}(\mathbf{x}_k | \mathbf{A}_{k-1}(\boldsymbol{\theta}_{k-1}) \mathbf{x}_{k-1}, \mathbf{Q}_{k-1}(\boldsymbol{\theta}_{k-1}))$$

$$p(\mathbf{y}_k | \mathbf{x}_k, \boldsymbol{\theta}_k) = \mathcal{N}(\mathbf{y}_k | \mathbf{H}_k(\boldsymbol{\theta}_k) \mathbf{x}_k, \mathbf{R}_k(\boldsymbol{\theta}_k))$$

$$p(\boldsymbol{\theta}_k | \boldsymbol{\theta}_{k-1}) = (\text{any given form}),$$

where

- \mathbf{x}_k is the state
- \mathbf{y}_k is the measurement
- $\boldsymbol{\theta}_k$ is an arbitrary latent variable
- **Given** the latent variables $\boldsymbol{\theta}_{1:T}$ the model is **Gaussian**.
- The RBPF uses SIR for the latent variables and computes the conditionally Gaussian part in closed form with Kalman filter.

- The **full posterior** at step k can be factored as

$$p(\theta_{0:k}, \mathbf{x}_{0:k} \mid \mathbf{y}_{1:k}) = p(\mathbf{x}_{0:k} \mid \theta_{0:k}, \mathbf{y}_{1:k}) p(\theta_{0:k} \mid \mathbf{y}_{1:k})$$

- The **first term** is Gaussian and computable with Kalman filter and RTS smoother
- For the **second term** we get the following recursion:

$$\begin{aligned} p(\theta_{0:k} \mid \mathbf{y}_{1:k}) &\propto p(\mathbf{y}_k \mid \theta_{0:k}, \mathbf{y}_{1:k-1}) p(\theta_{0:k} \mid \mathbf{y}_{1:k-1}) \\ &= p(\mathbf{y}_k \mid \theta_{0:k}, \mathbf{y}_{1:k-1}) p(\theta_k \mid \theta_{0:k-1}, \mathbf{y}_{1:k-1}) p(\theta_{0:k-1} \mid \mathbf{y}_{1:k-1}) \\ &= p(\mathbf{y}_k \mid \theta_{0:k}, \mathbf{y}_{1:k-1}) p(\theta_k \mid \theta_{k-1}) p(\theta_{0:k-1} \mid \mathbf{y}_{1:k-1}) \end{aligned}$$

- Let's take a look at the terms in

$$p(\mathbf{y}_k | \boldsymbol{\theta}_{0:k}, \mathbf{y}_{1:k-1}) p(\boldsymbol{\theta}_k | \boldsymbol{\theta}_{k-1}) p(\boldsymbol{\theta}_{0:k-1} | \mathbf{y}_{1:k-1})$$

- The **first term** can be computed by running Kalman filter with fixed $\boldsymbol{\theta}_{0:k}$ over the measurement sequence.
- The **second term** is just the dynamic model.
- The **third term** is the posterior from the previous step.

- We can form the **importance distribution** recursively:

$$\pi(\theta_{0:k} | \mathbf{y}_{1:k}) = \pi(\theta_k | \theta_{0:k-1}, \mathbf{y}_{1:k}) \pi(\theta_{0:k-1} | \mathbf{y}_{1:k-1})$$

- We then get the following **recursion** for the **weights**:

$$w_k^{(i)} \propto \frac{p(\mathbf{y}_k | \theta_{0:k-1}^{(i)}, \mathbf{y}_{1:k-1}) p(\theta_k^{(i)} | \theta_{k-1}^{(i)})}{\pi(\theta_k^{(i)} | \theta_{0:k-1}^{(i)}, \mathbf{y}_{1:k})} w_{k-1}^{(i)}$$

- Given the marginal posterior for $\theta_{0:k}$ we can recover the Gaussian part $\mathbf{x}_{0:k}$ with **Kalman filter and RTS smoother**.
- The **optimal importance distribution** takes the form

$$p(\theta_k | \mathbf{y}_{1:k}, \theta_{0:k-1}^{(i)}) \propto p(\mathbf{y}_k | \theta_k, \theta_{0:k-1}^{(i)}) p(\theta_k | \theta_{0:k-1}^{(i)}, \mathbf{y}_{1:k-1})$$

Rao-Blackwellized Particle Filter

- Perform Kalman filter predictions for each of the Kalman filter means and covariances in the particles $i = 1, \dots, N$ conditional on the previously drawn latent variable values $\theta_{k-1}^{(i)}$

$$\mathbf{m}_k^{-(i)} = \mathbf{A}_{k-1}(\theta_{k-1}^{(i)}) \mathbf{m}_{k-1}^{(i)}$$

$$\mathbf{P}_k^{-(i)} = \mathbf{A}_{k-1}(\theta_{k-1}^{(i)}) \mathbf{P}_{k-1}^{(i)} \mathbf{A}_{k-1}^T(\theta_{k-1}^{(i)}) + \mathbf{Q}_{k-1}(\theta_{k-1}^{(i)}).$$

- Draw new latent variables $\theta_k^{(i)}$ for each particle in $i = 1, \dots, N$ from the corresponding importance distributions

$$\theta_k^{(i)} \sim \pi(\theta_k \mid \theta_{0:k-1}^{(i)}, \mathbf{y}_{1:k}).$$

Rao-Blackwellized Particle Filter (cont.)

- Calculate new weights as follows:

$$w_k^{(i)} \propto w_{k-1}^{(i)} \frac{p(\mathbf{y}_k | \boldsymbol{\theta}_{0:k}^{(i)}, \mathbf{y}_{1:k-1}) p(\boldsymbol{\theta}_k^{(i)} | \boldsymbol{\theta}_{k-1}^{(i)})}{\pi(\boldsymbol{\theta}_k^{(i)} | \boldsymbol{\theta}_{0:k-1}^{(i)}, \mathbf{y}_{1:k})},$$

where the likelihood term is the marginal measurement likelihood of the Kalman filter:

$$\begin{aligned} p(\mathbf{y}_k | \boldsymbol{\theta}_{0:k}^{(i)}, \mathbf{y}_{1:k-1}) \\ = \mathcal{N} \left(\mathbf{y}_k \mid \mathbf{H}_k(\boldsymbol{\theta}_k^{(i)}) \mathbf{m}_k^{-(i)}, \mathbf{H}_k(\boldsymbol{\theta}_k^{(i)}) \mathbf{P}_k^{-(i)} \mathbf{H}_k^T(\boldsymbol{\theta}_k^{(i)}) + \mathbf{R}_k(\boldsymbol{\theta}_k^{(i)}) \right). \end{aligned}$$

- Then normalize the weights to sum to unity.

Rao-Blackwellized Particle Filter (cont.)

- Perform Kalman filter updates for each of the particles conditional on the drawn latent variables $\theta_k^{(i)}$

$$\mathbf{v}_k^{(i)} = \mathbf{y}_k - \mathbf{H}_k(\theta_k^{(i)}) \mathbf{m}_k^-$$

$$\mathbf{S}_k^{(i)} = \mathbf{H}_k(\theta_k^{(i)}) \mathbf{P}_k^{-(i)} \mathbf{H}_k^T(\theta_k^{(i)}) + \mathbf{R}_k(\theta_k^{(i)})$$

$$\mathbf{K}_k^{(i)} = \mathbf{P}_k^{-(i)} \mathbf{H}_k^T(\theta_k^{(i)}) \mathbf{S}_k^{(i)-1}$$

$$\mathbf{m}_k^{(i)} = \mathbf{m}_k^{-(i)} + \mathbf{K}_k^{(i)} \mathbf{v}_k^{(i)}$$

$$\mathbf{P}_k^{(i)} = \mathbf{P}_k^{-(i)} - \mathbf{K}_k^{(i)} \mathbf{S}_k^{(i)} [\mathbf{K}_k^{(i)}]^T.$$

- If the effective number of particles is too low, perform *resampling*.

Rao-Blackwellized Particle Filter: Properties

- The Rao-Blackwellized particle filter produces a set of weighted samples $\{w_k^{(i)}, \theta_k^{(i)}, \mathbf{m}_k^{(i)}, \mathbf{P}_k^{(i)} : i = 1, \dots, N\}$
- The **expectation of a function** $\mathbf{g}(\cdot)$ can be approximated as

$$\mathbb{E}[\mathbf{g}(\mathbf{x}_k, \theta_k) | \mathbf{y}_{1:k}] \approx \sum_{i=1}^N w_k^{(i)} \int \mathbf{g}(\mathbf{x}_k, \theta_k^{(i)}) \mathcal{N}(\mathbf{x}_k | \mathbf{m}_k^{(i)}, \mathbf{P}_k^{(i)}) d\mathbf{x}_k.$$

- Approximation of the **filtering distribution** is

$$p(\mathbf{x}_k, \theta_k | \mathbf{y}_{1:k}) \approx \sum_{i=1}^N w_k^{(i)} \delta(\theta_k - \theta_k^{(i)}) \mathcal{N}(\mathbf{x}_k | \mathbf{m}_k^{(i)}, \mathbf{P}_k^{(i)}).$$

- It is possible to do **approximate Rao-Blackwellization** by replacing the Kalman filter with a Gaussian filter.

Rao-Blackwellization of Static Parameters

- Rao-Blackwellization can sometimes be used in models of the form

$$\begin{aligned}\mathbf{x}_k &\sim p(\mathbf{x}_k \mid \mathbf{x}_{k-1}, \theta) \\ \mathbf{y}_k &\sim p(\mathbf{y}_k \mid \mathbf{x}_k, \theta) \\ \theta &\sim p(\theta),\end{aligned}$$

where vector θ contains the **unknown static parameters**.

- Possible if the posterior distribution of parameters θ depends only on some **sufficient statistics** T_k :

$$\mathbf{T}_k = \mathbf{T}_k(\mathbf{x}_{1:k}, \mathbf{y}_{1:k})$$

- We also need to have a recursion rule for the **sufficient statistics**.
- Can be extended to time-varying parameters.

Particle Filter: Advantages

- No restrictions in model – can be applied to **non-Gaussian models, hierarchical models etc.**
- **Global** approximation.
- Approaches the **exact solution**, when the number of samples goes to infinity.
- In its **basic** form, very **easy to implement**.
- Superset of other filtering methods – Kalman filter is a Rao-Blackwellized particle filter with one particle.

Particle Filter: Disadvantages

- Computational requirements much higher than of the Kalman filters.
- Problems with nearly noise-free models, especially with accurate dynamic models.
- Good importance distributions and efficient Rao-Blackwellized filters quite tricky to implement.
- Very hard to find programming errors (i.e., to debug).

- **Particle filters** use **weighted set of samples** (particles) for approximating the filtering distributions.
- **Sequential importance resampling (SIR)** is the general framework and **bootstrap filter** is a simple special case of it.
- **EKF, UKF and other Gaussian filters** can be used for forming good **importance distributions**.
- In **Rao-Blackwellized particle filters** a part of the state is sampled and part is integrated in closed form with Kalman filter.

- The discretized pendulum model:

$$\begin{pmatrix} x_k^1 \\ x_k^2 \end{pmatrix} = \underbrace{\begin{pmatrix} x_{k-1}^1 + x_{k-1}^2 \Delta t \\ x_{k-1}^2 - g \sin(x_{k-1}^1) \Delta t \end{pmatrix}}_{\mathbf{f}(\mathbf{x}_{k-1})} + \begin{pmatrix} 0 \\ q_{k-1} \end{pmatrix}$$
$$y_k = \underbrace{\sin(x_k^1)}_{\mathbf{h}(\mathbf{x}_k)} + r_k,$$

- \Rightarrow *Matlab demonstration*